

Search Trees

Data Structures and Algorithms

Nanjing University, Fall 2021

郑朝栋

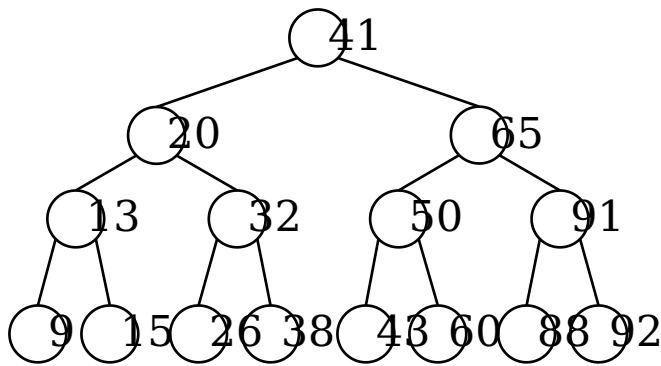
Efficient implementation of OSet

	Search(S,k)	Insert(S,x)	Remove(S,x)
BinarySearchTree	$O(h)$ worst-case	$O(h)$ worst-case	$O(h)$ worst-case
Treap	$O(\log n)$ in expectation	$O(\log n)$ in expectation	$O(\log n)$ in expectation

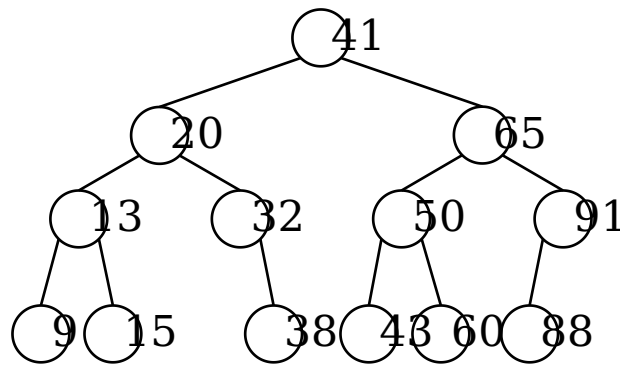
A data structure supporting OSet operations in $O(\log n)$ time, even in worst-case?

“Balanced” BST

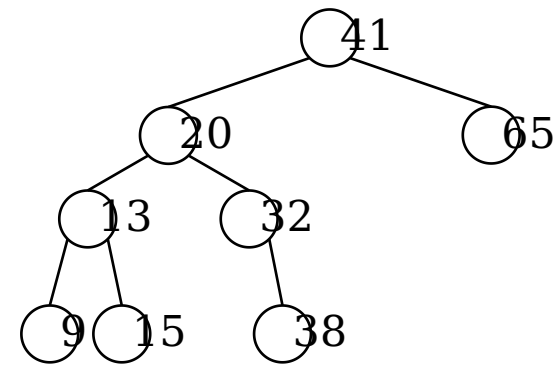
- What does it mean to be “balanced”?
 - **Perfectly Balanced.** (For each node, two subtrees have same height.)
 - **Almost Perfectly Balanced.**
 - **Not Perfectly Balanced.**
- An n -node BST is “balanced” if it has height $O(\log n)$.



Perfectly Balanced



Almost Perfectly Balanced



Not Perfectly Balanced

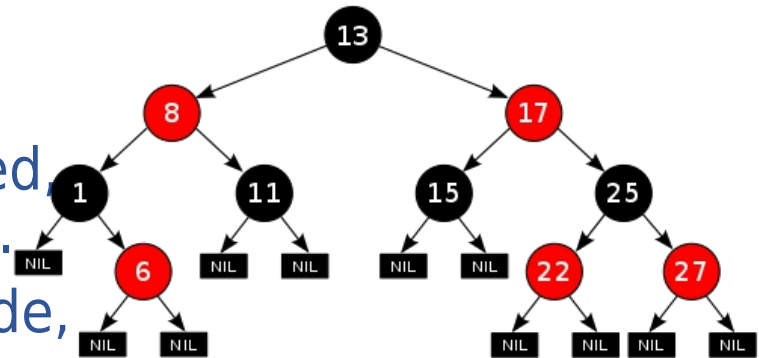
“Balanced” ~~Binary~~ Search Trees

- AVL tree (Adelson-Velsii & Landis, 1962)
- B-tree (Bayer & McCreight, 1970)
- Red-black tree (Bayer, 1972)
- Splay tree (Sleator & Tarjan, 1985)
- Treap (Seidel & Aragon, 1996)
- Skip list (Pugh, 1989)
- and so on ...

Red-black Tree (RB-Tree)

A **Red-black Tree (RB-Tree)** is a BST in which each node has a *color*, and satisfies the following properties:

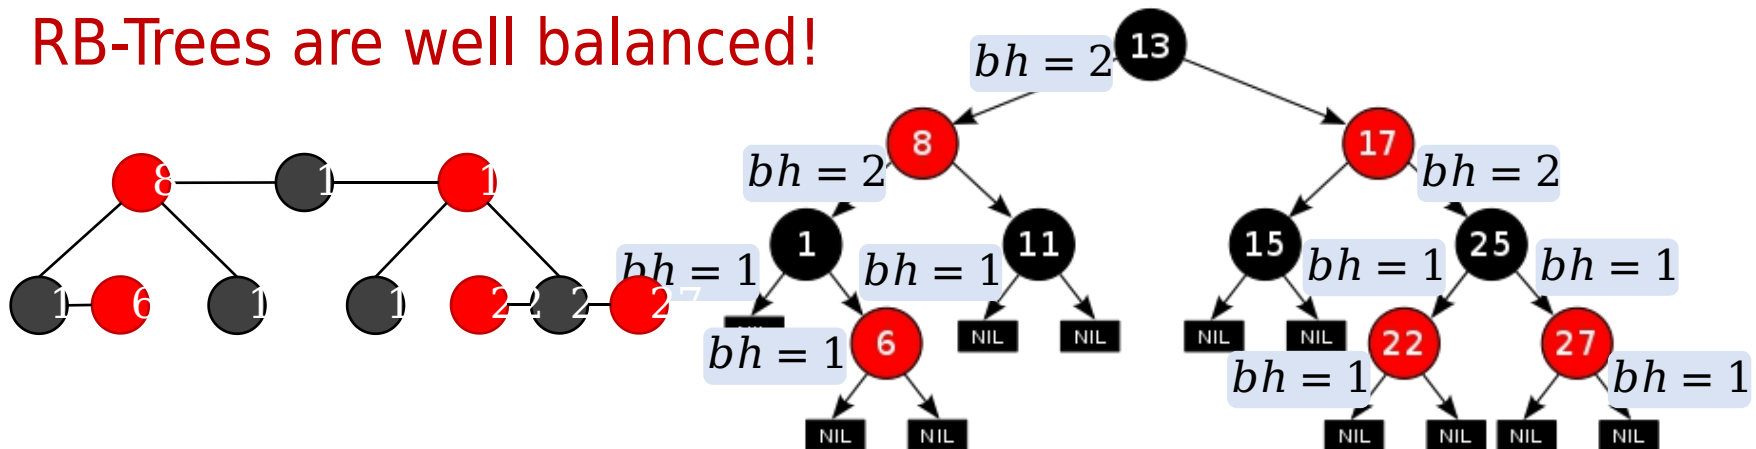
- Every node is either red or black.
- The root is black.
- Every leaf is black.
- **[no-red-edge]** If a node is red, then both its children are black.
- **[black-height]** For every node, all paths from the node to its descendant leaves contain same number of black nodes.



Black Height

- Call the number of black nodes on any simple path from, but not including, a node x down to a leaf the **black-height** of the node, denoted by $bh(x)$.
- Due to black-height property, from the black-height perspective, RB-Trees are “perfectly balanced”.
- Due to no-red-edge property, actual height of a RB-Tree does not deviate a lot from its black-height.

RB-Trees are well balanced!



Height of RB-Trees

- **Claim:** In a RB-Tree, the subtree rooted at x contains at least $2^{bh(x)} - 1$ internal nodes.
- **Proof** (via induction on height of x):
- **[Basis]** If x is a leaf, $bh(x) = 0$ and the claim holds.
- **[Hypothesis]** The claim holds for all nodes with height at most $h - 1$.
- **[Inductive Step]** Consider a node x with height $h \geq 1$. It must have two children. (Why?) So the number of internal nodes rooted at x is:
$$\begin{aligned} &\geq 1 + (2^{bh(x.left)} - 1) + (2^{bh(x.right)} - 1) \\ &\geq 1 + (2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) \\ &= 2^{bh(x)} - 1 \end{aligned}$$

Height of RB-Trees

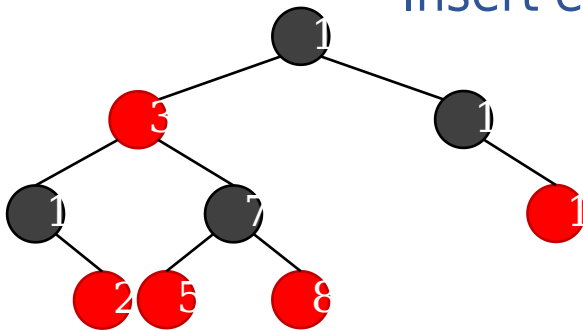
- **Claim:** In a RB-Tree, the subtree rooted at x contains at least $2^{bh(x)} - 1$ internal nodes.
- Due to no-red-edge: $h = height(root) \leq 2 \cdot bh(root)$.
- $n \geq 2^{bh(root)} - 1 \geq 2^{h/2} - 1$, implying $h \leq 2 \cdot \lg(n + 1)$.
- **Theorem:** The height of an n -node RB-Tree is $O(\log n)$.
- RB-Trees support Search, Min, Max, Predecessor, Successor operations in worst-case $O(\log n)$ time!
- What about Insert and Remove?!

Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
 - **Step 2:** Fix any violated properties.
 - No fix is needed if z has a black parent.
- Maintain black-height, fix no-red-edge if necessary.

Example:

Insert element with key 2.



RB-Tree Properties:

- Each node is red or black.
- Root is black.
- Leaves are black.
- No-red-edge property.
- Black-height property.

Insert node z into an RB-Tree

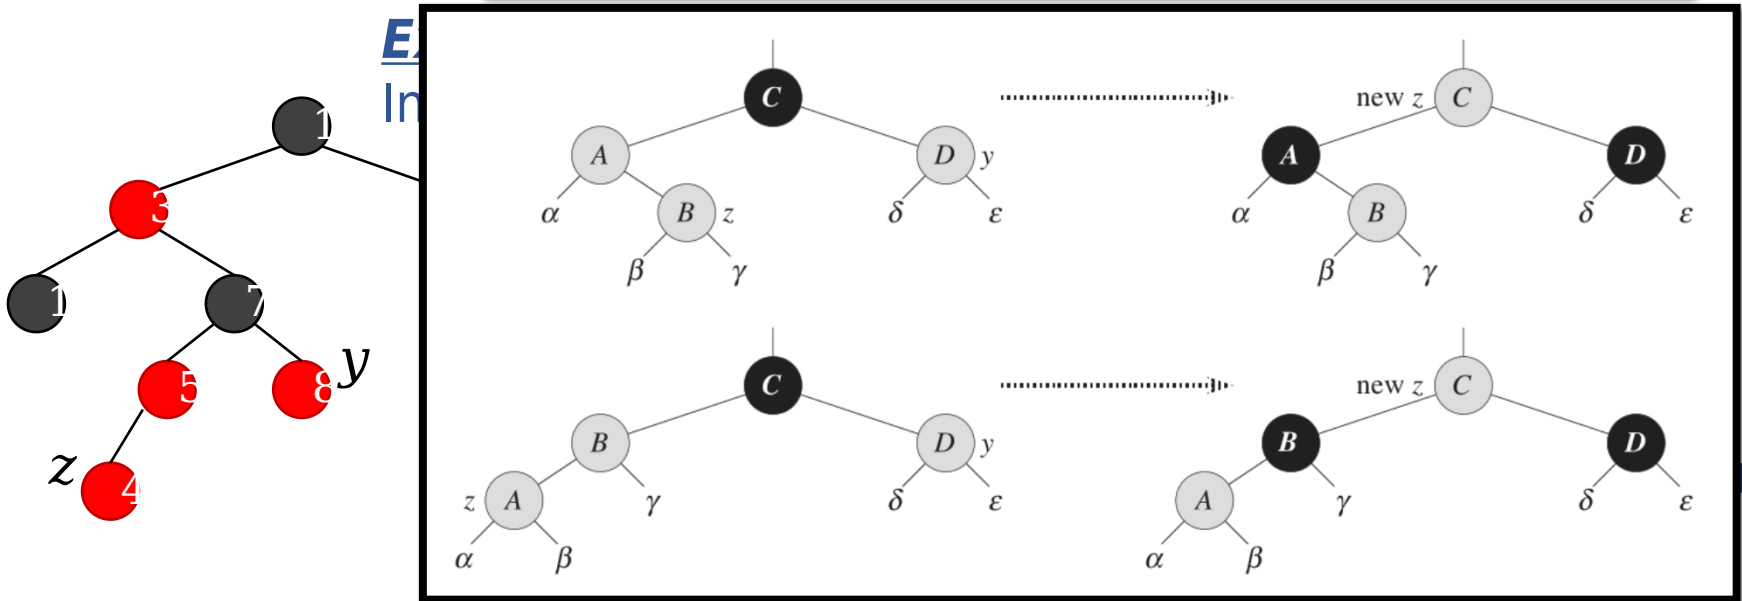
- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **Case 0:** z becomes the root of the RB-Tree.
 - **Fix:** simply recolor z to be black.

RB-Tree Properties:

- Each node is red or black.
- Root is black (**easy fix**).
- Leaves are black.
- No-red-edge property (**fix**).
- Black-height property (**maintain**).

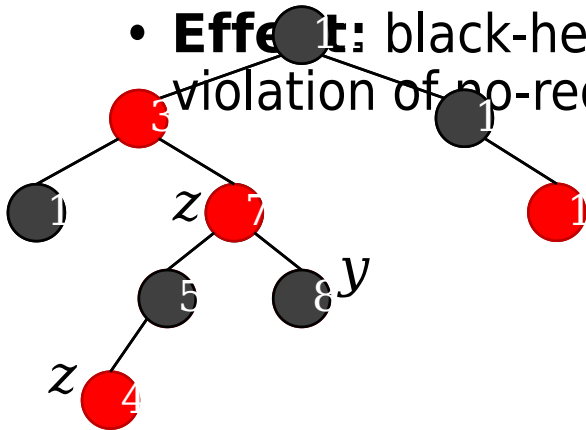
Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **Case 1:** z 's Black-height property maintained, and we as red uncle y "push-up" violation of no-red-edge property



Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **Case 1:** z 's parent is red (so z has black grandparent), and has red uncle y .
 - **Fix:** recolor z 's parent and uncle to black, recolor z 's grandparent to red.
 - **Effect:** black-height property maintained, and we "push-up" violation of no-red-edge property.

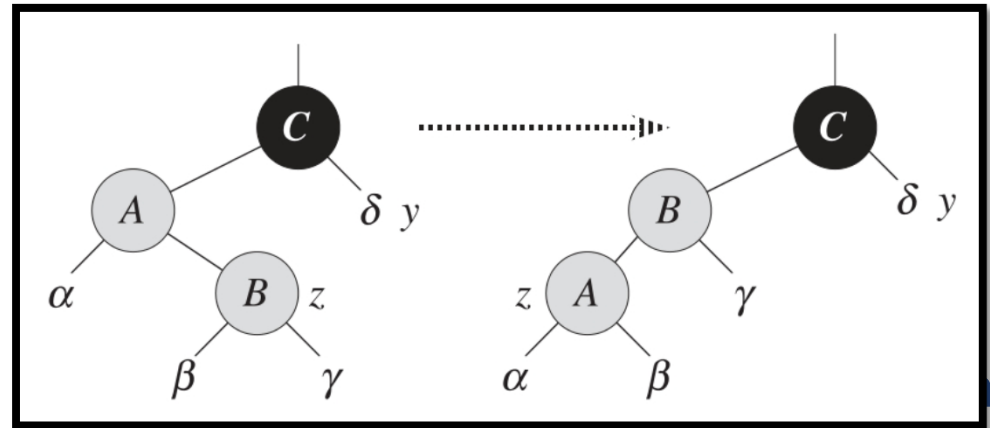
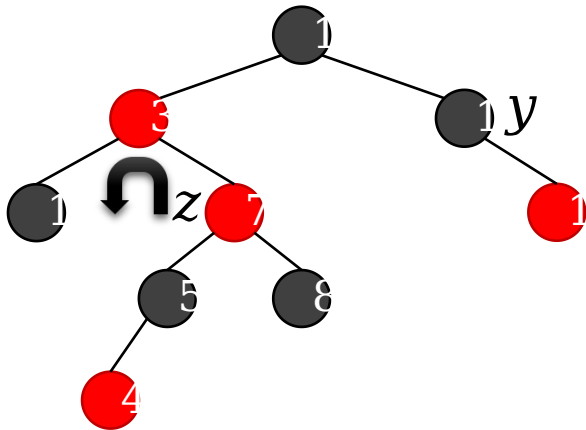


RB-Tree Properties:

- Each node is red or black.
- Root is black (**easy fix**).
- Leaves are black.
- No-red-edge property (**fix**).
- Black-height property (**maintain**).

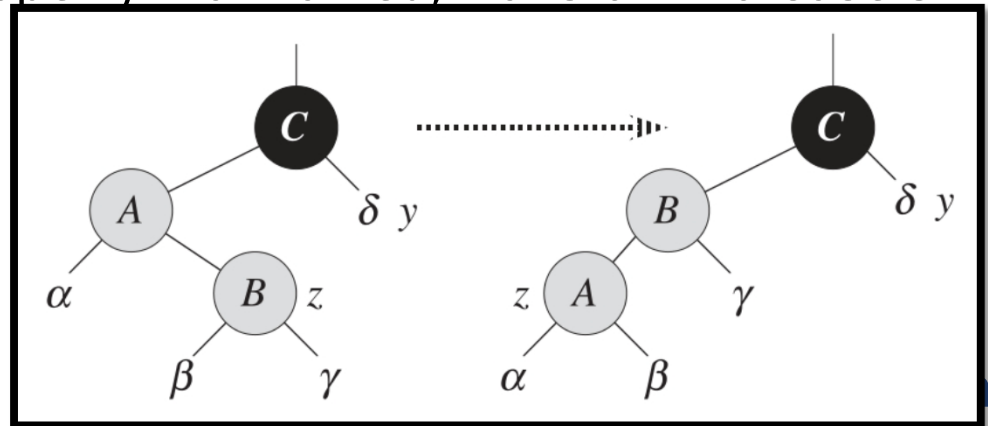
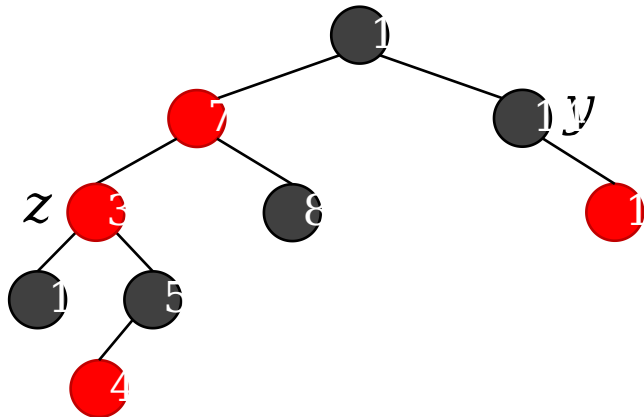
Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **Case 2:** z 's parent is red, has black uncle y . z is right child of its parent.
 - **Fix:** "left-rotate" at z 's parent.



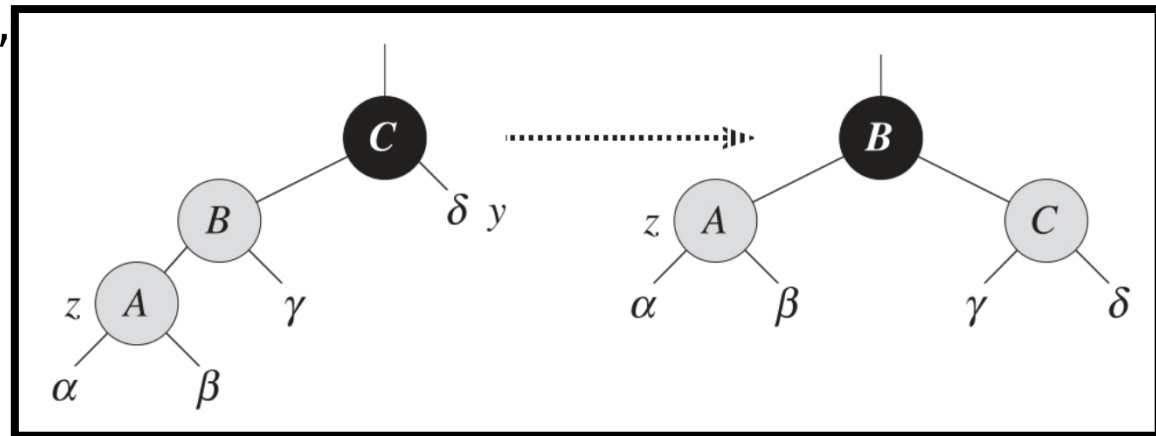
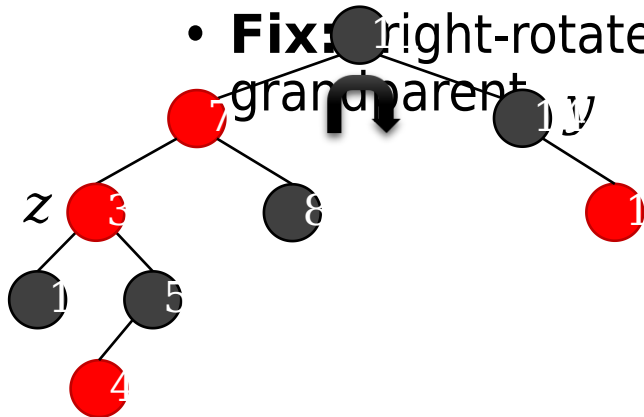
Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **Case 2:** z 's parent is red, has black uncle y . z is right child of its parent.
 - **Fix:** "left-rotate" at z 's parent.
 - **Effect:** black-height property maintained, transform to Case 3.



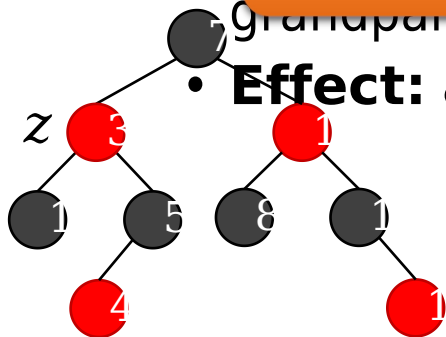
Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **Case 2:** z 's parent is red, has black uncle y . z is right child of its parent.
 - **Case 3:** z 's parent is red, has black uncle y . z is left child of its parent.
 - **Fix:** "right-rotate" grandparent

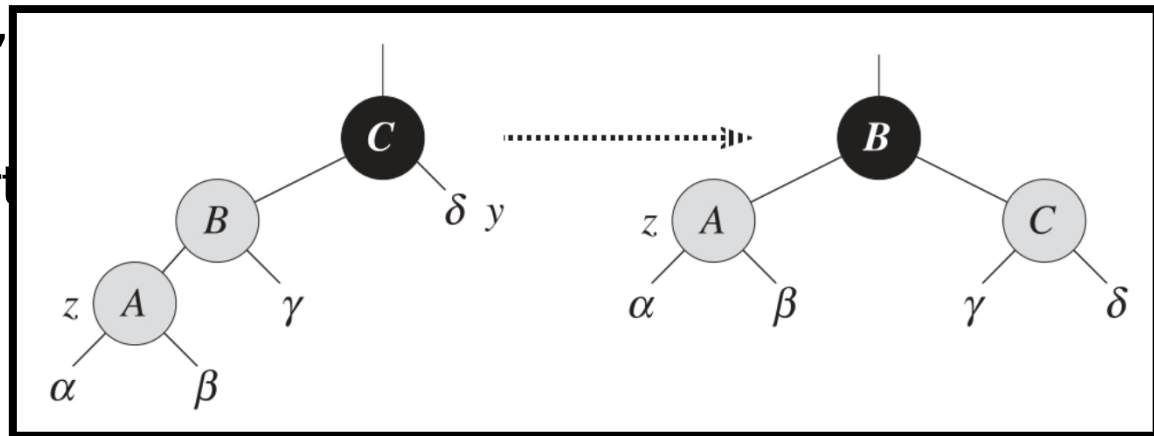


Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **Case 2:** z 's parent is red, has black uncle y . z is right child of its parent.
 - **Case 3:** z 's parent is red, has black uncle y . z is left child of its parent.
 - **Effect:** all properties are restored.

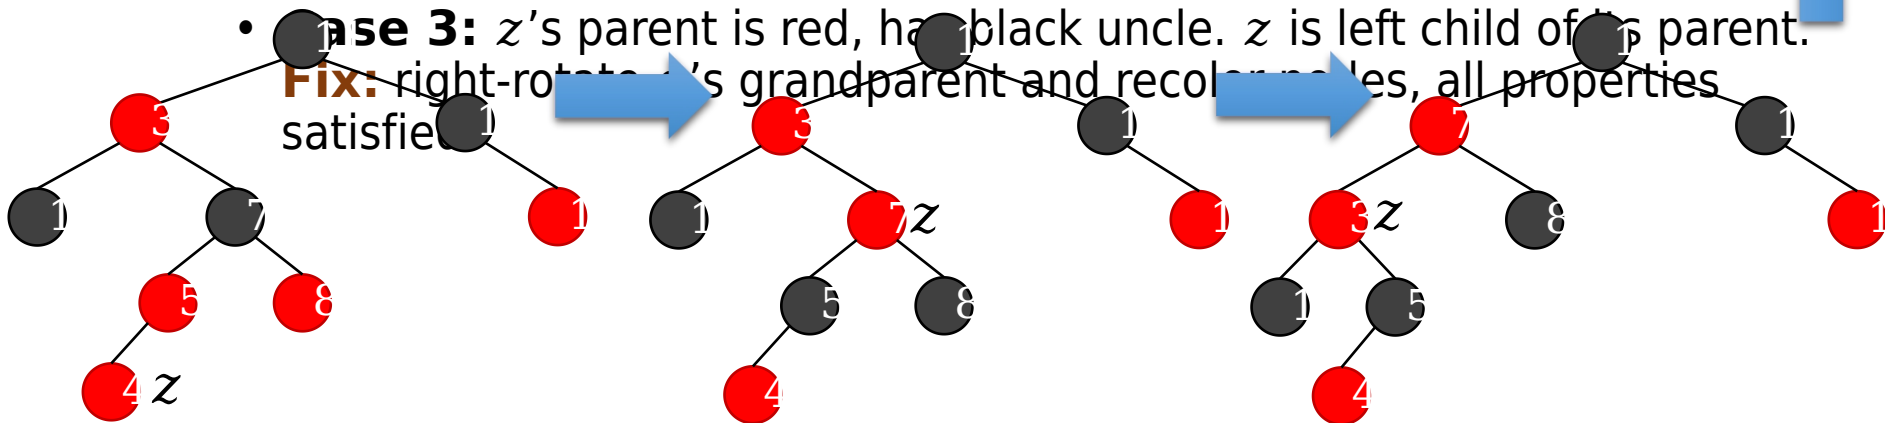
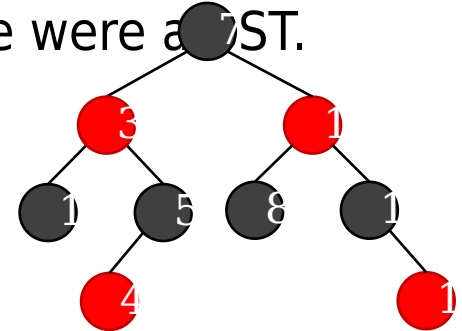


We are done!"



Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **No-Fix-Needed:** z has a black parent.
 - **Case 0:** z becomes the root.
Fix: recolor z to be black.
 - **Case 1:** z 's parent is red, has red uncle.
Fix: recoloring to push-up "no-red-edge" violation (maintain "black-height").
 - **Case 2:** z 's parent is red, has black uncle. z is right child of its parent.
Fix: left-rotate z 's parent to transform to Case 3 (maintain "black-height").
 - **Case 3:** z 's parent is red, has black uncle. z is left child of its parent.
Fix: right-rotate z 's grandparent and recolor nodes, all properties satisfied.



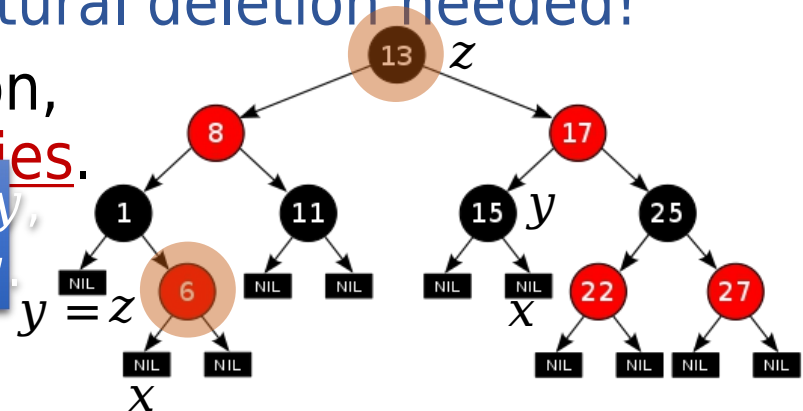
Insert node z into an RB-Tree

- **Step 1:** Color z as red and insert as if the RB-tree were a BST.
- **Step 2:** Fix any violated properties.
 - **No-Fix-Needed:** z has a black parent.
 - **Case 0:** z becomes the root.
Fix: recolor z to be black.
 - **Case 1:** z 's parent is red, has red uncle.
Fix: recoloring to push-up “no-red-edge” violation (maintain “black-height”).
 - **Case 2:** z 's parent is red, has black uncle. z is right child of its parent.
Fix: left-rotate z 's parent to transform to Case 3 (maintain “black-height”).
 - **Case 3:** z 's parent is red, has black uncle. z is left child of its parent.
Fix: right-rotate z 's grandparent and recolor nodes, all properties satisfied.
- Time Complexity of **Insert** operation?
 - $O(h) = O(\log n)$ time. (Case 1 appears at most $O(h)$ times.)
 - $O(1)$ rotations. (**Insert** has limited impact on tree shape.)

Remove node z from an RB-Tree

- If z 's right child is an external node, then z is the node to be deleted **structurally**: subtree rooted at $z.left$ will replace z .
- If z 's right child is an internal node, then let y be the min node in subtree rooted at $z.right$. Overwrite z 's info with y 's info, and y is the node to be deleted **structurally**: subtree rooted at $y.right$ will replace y .
- Either way, only **one** structural deletion needed!
- Apply the structural deletion, and **repair violated properties**.

Call the node to be deleted structurally y , and let x be the node that will replace y .

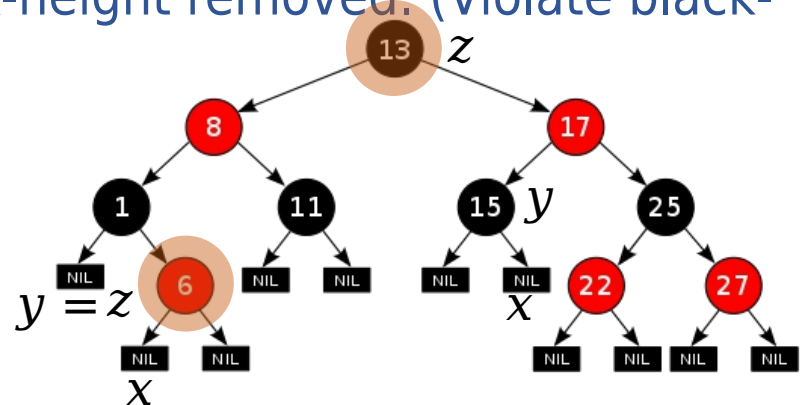


Remove node z from an RB-Tree

- **Step 1:** Identify the structural deletion.
- **Step 2:** Apply the structural deletion. (Maintain BST property.)
- **Step 3:** Repair violated RB-tree properties. (Maintain BST property.)
 - If y is a **red** node: no violations.
 - If y is a **black** node and x is a **red** node: recolor x to black and done.
 - If y is a **black** node and x is a **black** node:
 - y 's contribution to black-height removed. (Violate black-height.)

RB-Tree Properties:

- Each node is red or black.
- Root is black.
- Leaves are black.
- No-red-edge property.
- Black-height property.

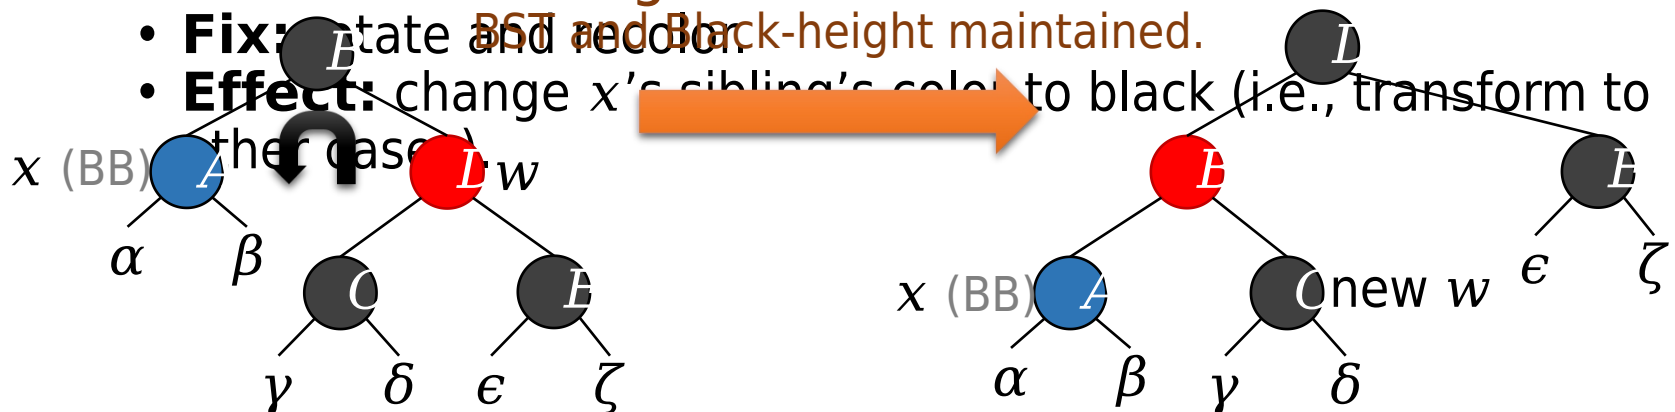


Remove node z from an RB-Tree

- **Step 1&2:** Find & apply structural deletion. (Maintain BST property.)
 - Let y be the structurally removed node, and x takes its place.
- **Step 3:** Repair violated RB-tree properties. (Maintain BST property.)
 - Assume black x is left child of its parent after taking black y 's place.
 - Focus on fixing black-height property.
- **Case 1:** x 's sibling w is red.

• **Fix:** rotate BST and recolor. (BST and Black-height maintained.)

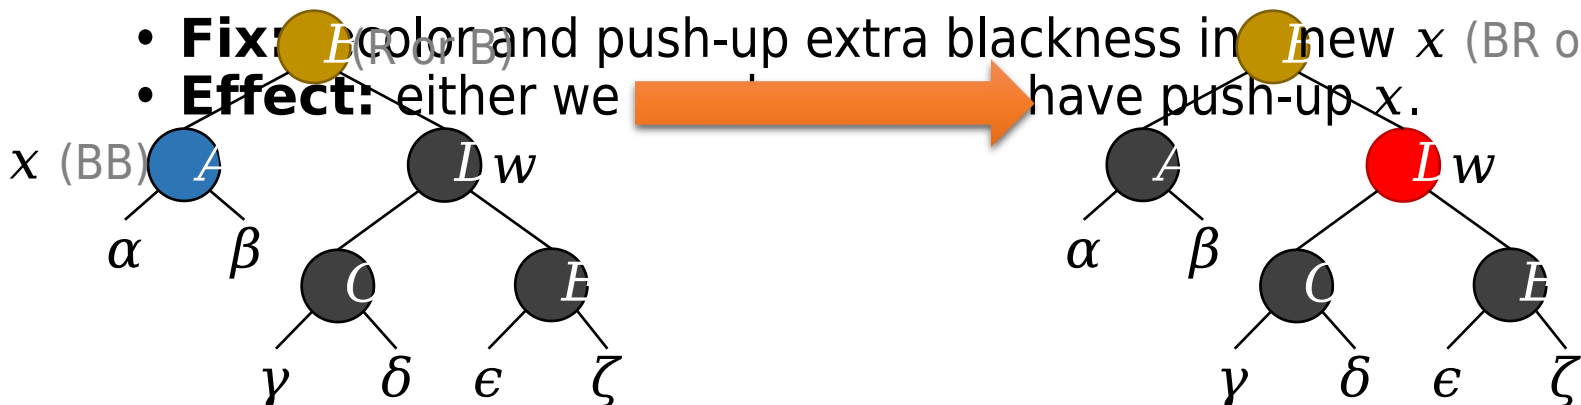
• **Effect:** change x 's sibling's color to black (i.e., transform to other case).



Remove node z from an RB-Tree

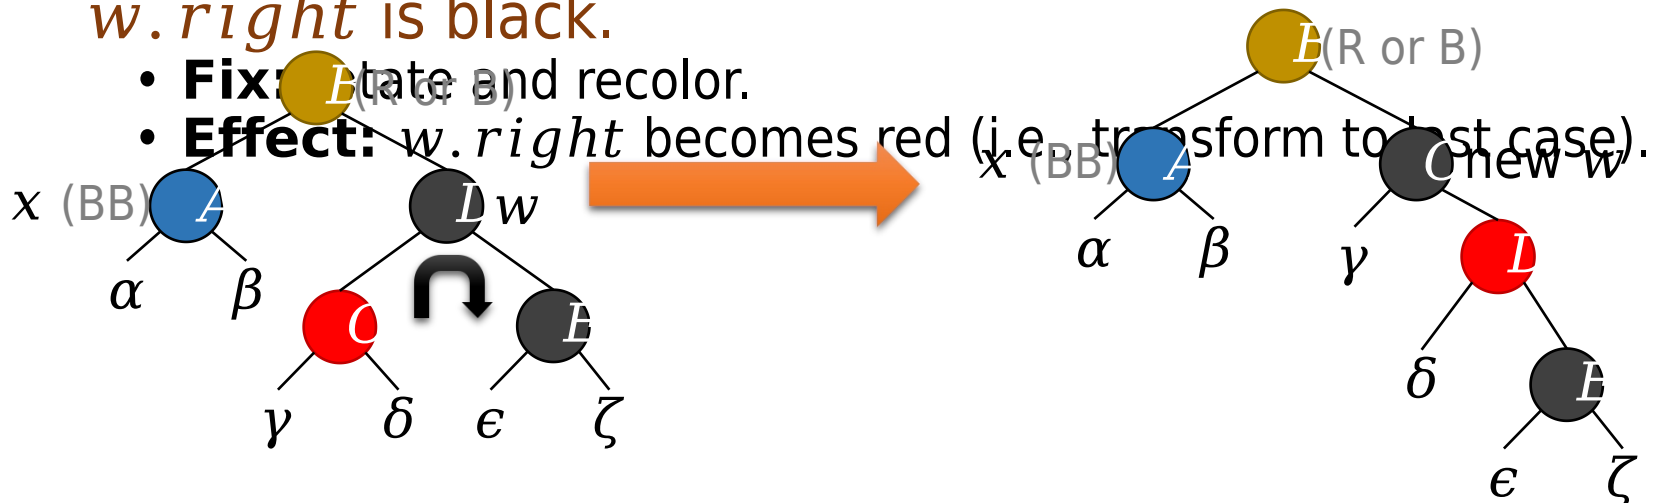
- **Step 1&2:** Find & apply structural deletion. (Maintain BST property.)
 - Let y be the structurally removed node, and x takes its place.
- **Step 3:** Repair violated RB-tree properties. (Maintain BST property.)
 - Assume x is left child of its parent.
 - Focus on fixing black-height property.
- **Case 2:** x 's sibling w is black, and both w 's children are black.

- **Fix:** Recolor and push-up extra blackness in new x (BR or BB)
- **Effect:** either we have push-up x .



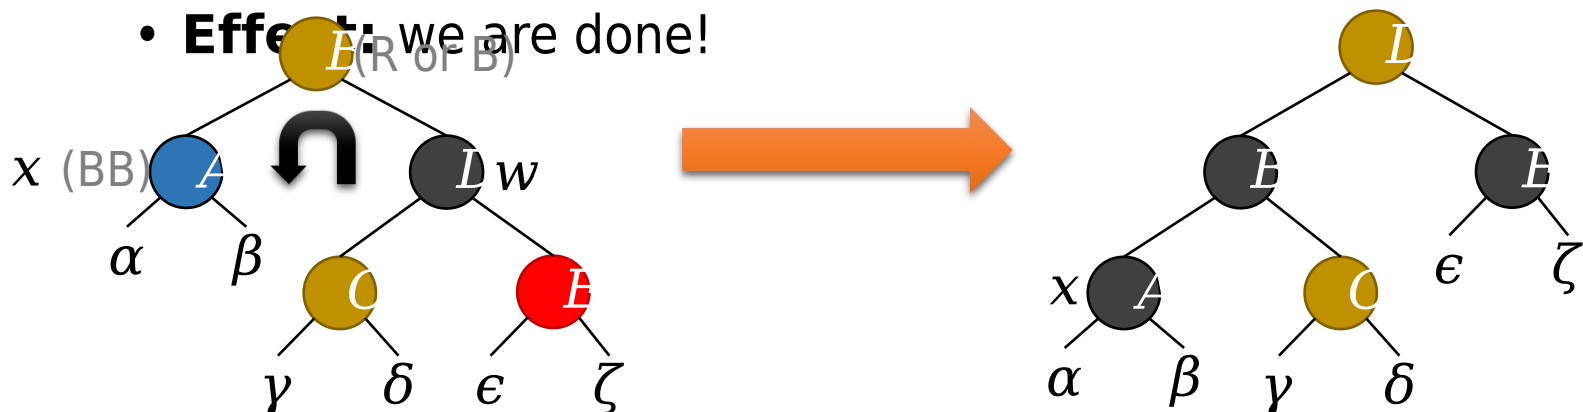
Remove node z from an RB-Tree

- **Step 1&2:** Find & apply structural deletion. (Maintain BST property.)
 - Let y be the structurally removed node, and x takes its place.
- **Step 3:** Repair violated RB-tree properties. (Maintain BST property.)
 - Assume x is left child of its parent.
 - Focus on fixing black-height property.
- **Case 3:** x 's sibling w is black, $w.left$ is red and $w.right$ is black.



Remove node z from an RB-Tree

- **Step 1&2:** Find & apply structural deletion. (Maintain BST property.)
 - Let y be the structurally removed node, and x takes its place.
- **Step 3:** Repair violated RB-tree properties. (Maintain BST property.)
 - Assume x is left child of its parent.
 - Focus on fixing black-height property.
- **Case 4:** x 's sibling w is black, $w.right$ is red.
 - **Fix:** rotate and recolor.
 - **Effect:** we are done!



Remove node z from an RB-Tree

- **Step 1&2:** Find & apply structural deletion (maintain BST property.)
 - Let y be the structurally removed node, and x be the node that replaces y .
 - $O(h)$
 - $= O(\log n)$
- **Step 3:** Repair violated RB-tree properties. (Maintain BST property.)
 - If x is not double-black: then done or easy fix.
 - If x is double-black:
 - **Case 1:** rotate and recolor; transform to other cases.
 - **Case 2:** recolor; done or push-up violations.
 - **Case 3:** rotate and recolor; transform to Case 4.
 - **Case 4:** rotate and recolor; then done.
- Time complexity of **Remove** operation?
 - $O(h) = O(\log n)$ time.
(For each case, in $O(1)$ time, either done or push-up violation.)
 - $O(1)$ rotations. [**Remove** has limited impact on tree shape.]
(Entering Case 2 from Case 1 will finish the fixing process.)

Efficient implementation of **OSet**

	Search(S,k)	Insert(S,x)	Remove(S,x)
BinarySearchTree	$O(h)$ worst-case	$O(h)$ worst-case	$O(h)$ worst-case
Treap	$O(\log n)$ in expectation	$O(\log n)$ in expectation	$O(\log n)$ in expectation
RB-Tree	$O(\log n)$ worst-case	$O(\log n)$ worst-case	$O(\log n)$ worst-case

Efficiency versus Simplicity

“Balanced” ~~Binary~~ Search Trees

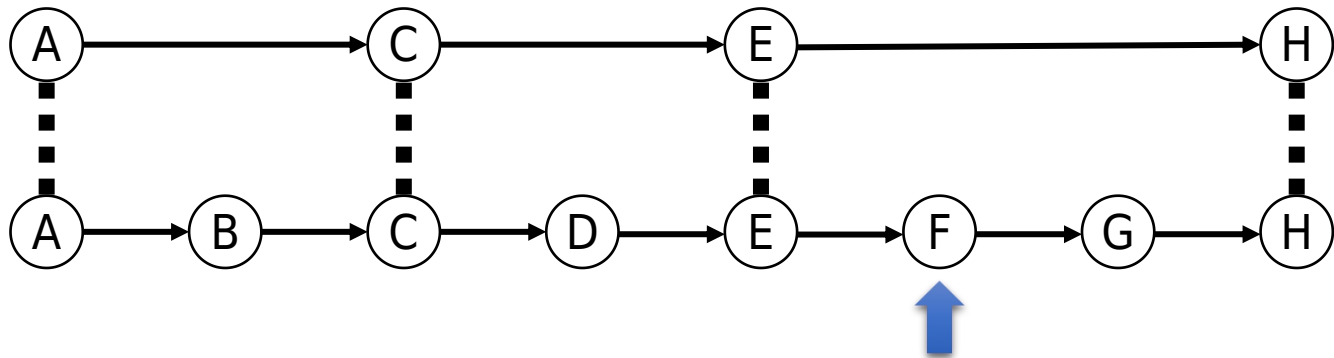
- AVL tree (Adelson-Velsii & Landis, 1962)
- B-tree (Bayer & McCreight, 1970)
- Red-black tree (Bayer, 1972)
- Splay tree (Sleator & Tarjan, 1985)
- Treap (Seidel & Aragon, 1996)
- Skip list (Pugh, 1989)
- and so on ...

SkipList

	Search(S,k)	Insert(S,x)	Remove(S,x)
SortedLinkedList	$O(n)$	$O(n)$	$O(1)$

Q: Why sorted linked-list is slow?

A: To reach an element, you have to move from current position to destination **one element at a time**.



SkipList

	Search(S,k)	Insert(S,x)	Remove(S,x)
SortedLinkedList	$O(n)$	$O(n)$	$O(1)$

Q: Why sorted linked-list is slow?

A: To reach an element, you have to move from current position to destination **one element at a time**.

Let's build an "expressway" Search cost is reduced by half!

Why stop at one layer of "expressway"?!

Example: search for 8.

SkipList

Search can be done in $O(\log n)$ time

Space complexity $\approx 2n$

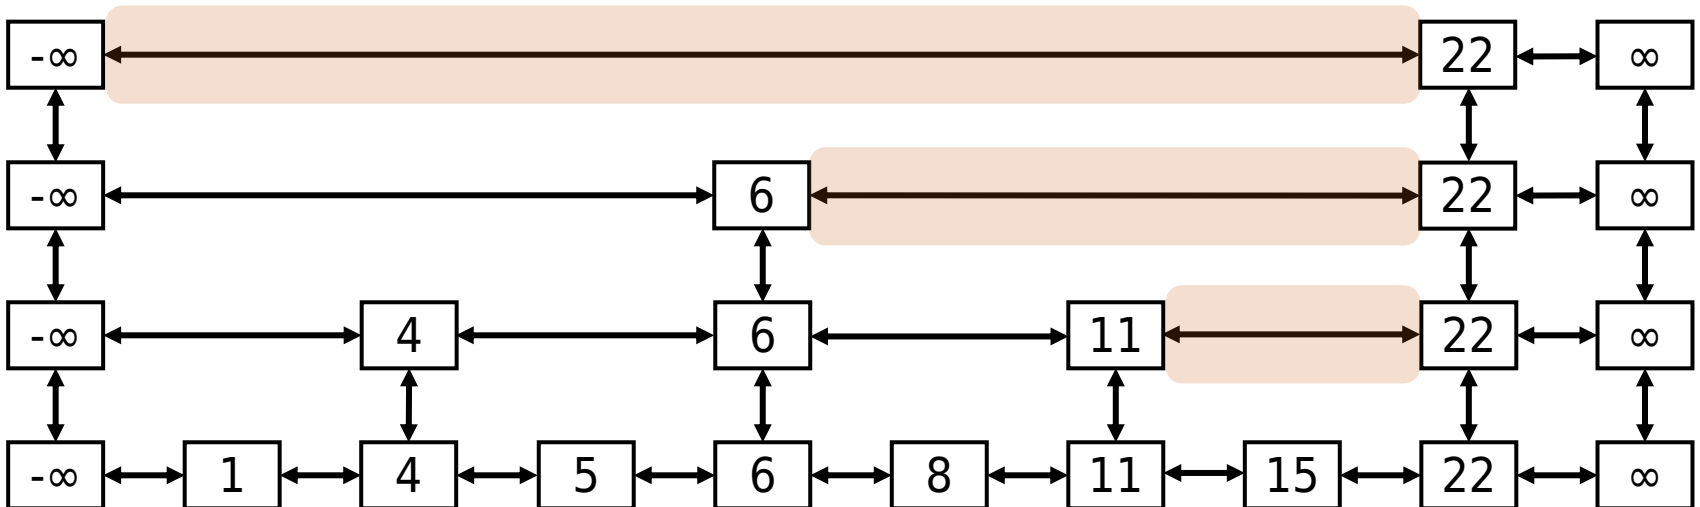
Build multiple “expressways”:

Reduce number of elements by half at each level.

This is just binary search: reduce search range by half at each level.

This is very efficient: spend $O(1)$ time at each level, and $O(\log n)$ levels

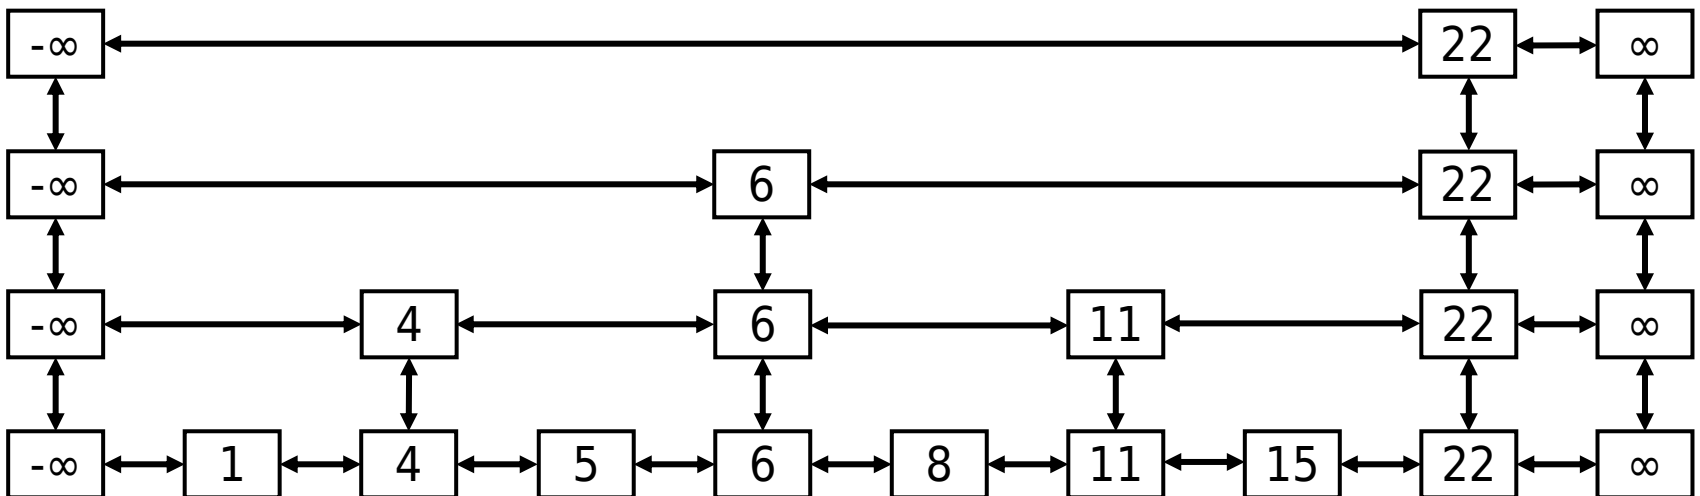
Example One: search for 15. Example Two: search for 14.



SkipList

	Search(S,k)	Insert(S,x)	Remove(S,x)
SortedLinkedList	$O(n)$	$O(n)$	$O(1)$
Static-SkipList	$O(\log n)$???	???

Efficient **Search** with limited space overhead.
But how to implement **Insert** and **Remove**?



The real **SkipList**

Insert(L,x):

level = 1, done = false

while (!done)

 Insert x into level k list.

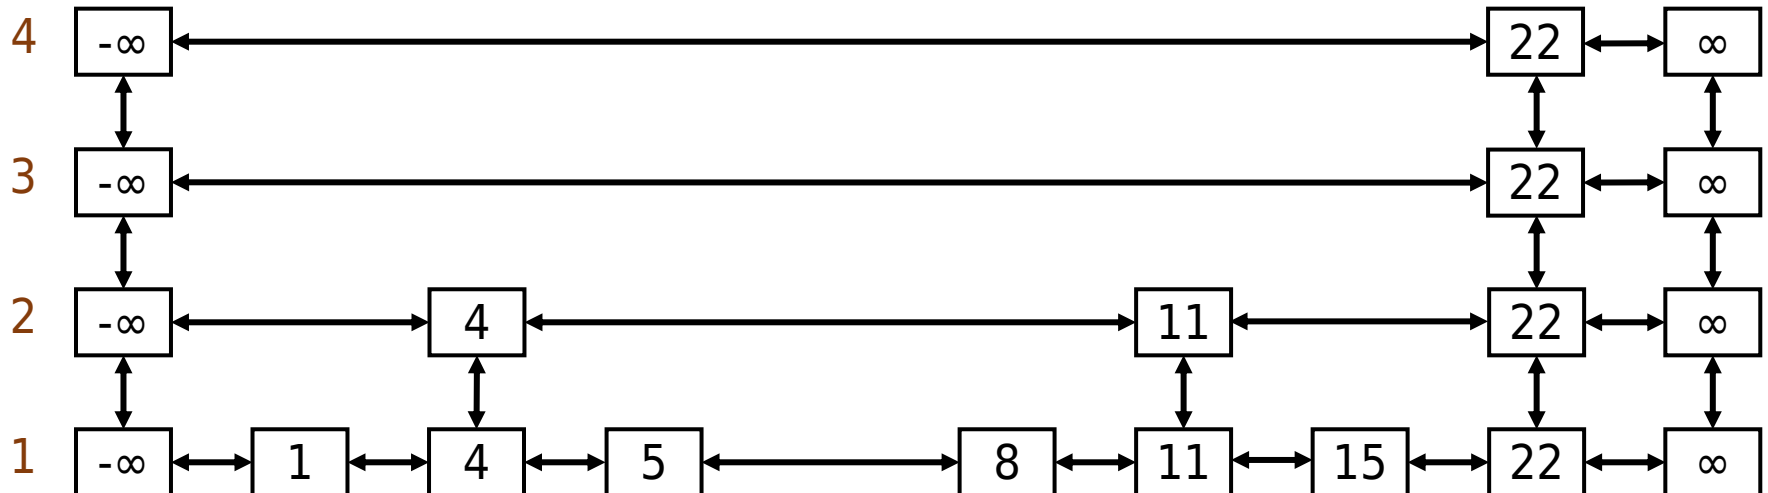
 Flip a fair coin:

 With probability 1/2: done = true

 With probability 1/2: k = k+1

Example: insert 7.

Lvl:



The real **SkipList**

Insert(L,x):

level = 1, done = false

while (!done)

 Insert x into level k list.

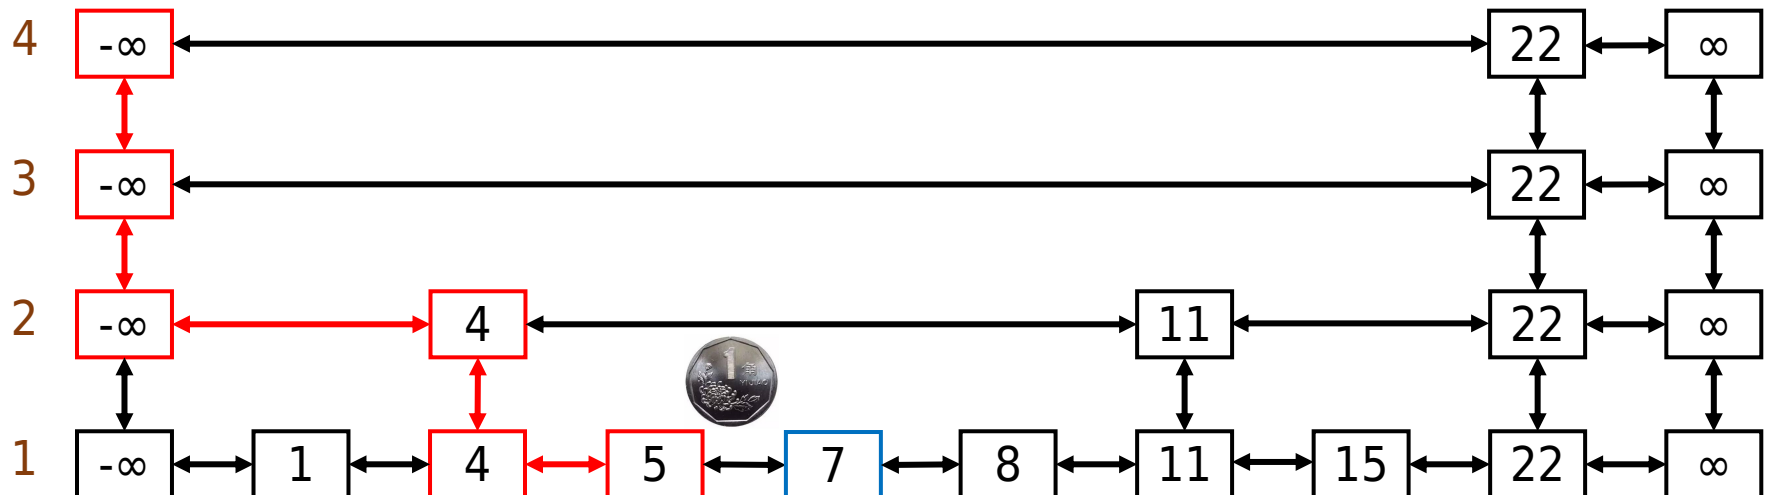
 Flip a fair coin:

 With probability 1/2: done = true

 With probability 1/2: k = k+1

Example: insert 7.

Lvl:



The real SkipList

Insert(L,x):

level = 1, done = false

while (!done)

 Insert x into level k list.

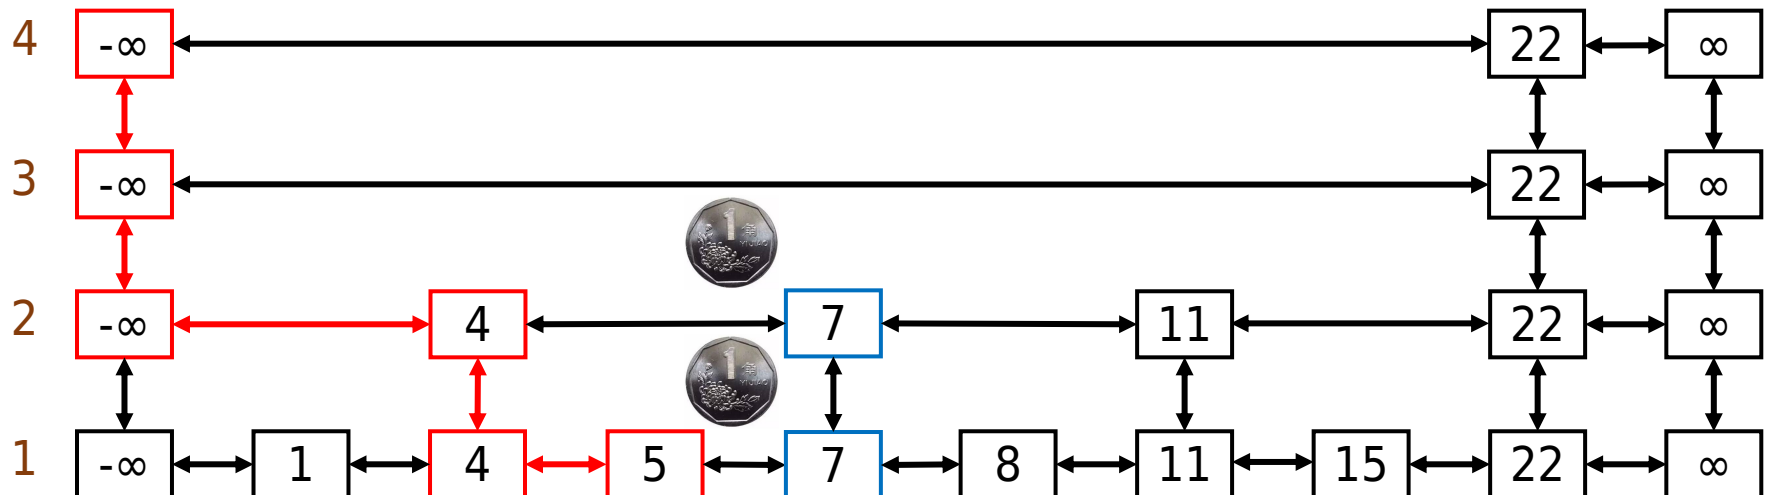
 Flip a fair coin:

 With probability 1/2: done = true

 With probability 1/2: k = k+1

Example: insert 7.

Lvl:



The real SkipList

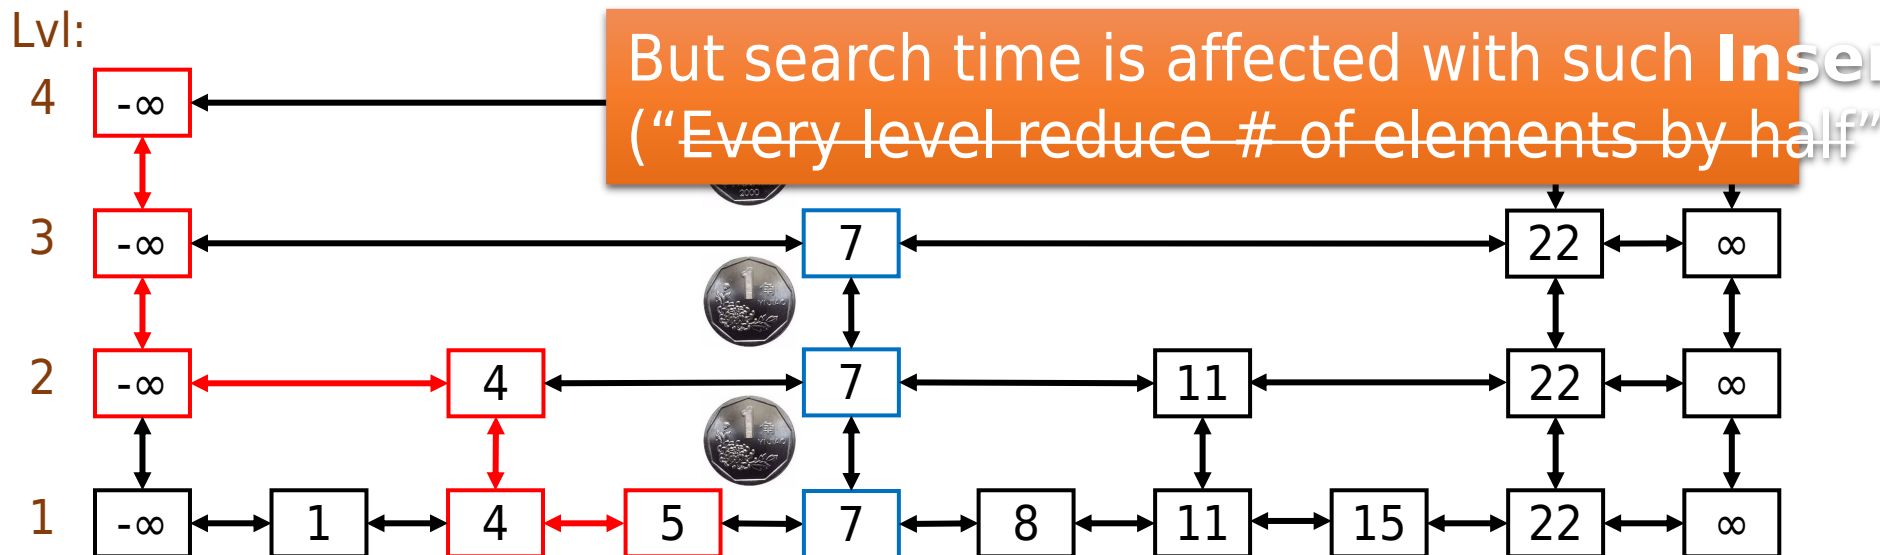
Insert(L,x):

```
level = 1, done = false
while (!done)
  Insert x into level k list.
  Flip a fair coin:
    With probability 1/2: done = true
    With probability 1/2: k = k+1
```

Time complexity of **Insert**:

- $O(1)$ in expectation.
- $O(\log n)$ with high probability.
(with prob. $\geq 1 - 1/n^{\Theta(1)}$)

Max level of n -element **SkipList** is $O(\log n)$ with high probability.



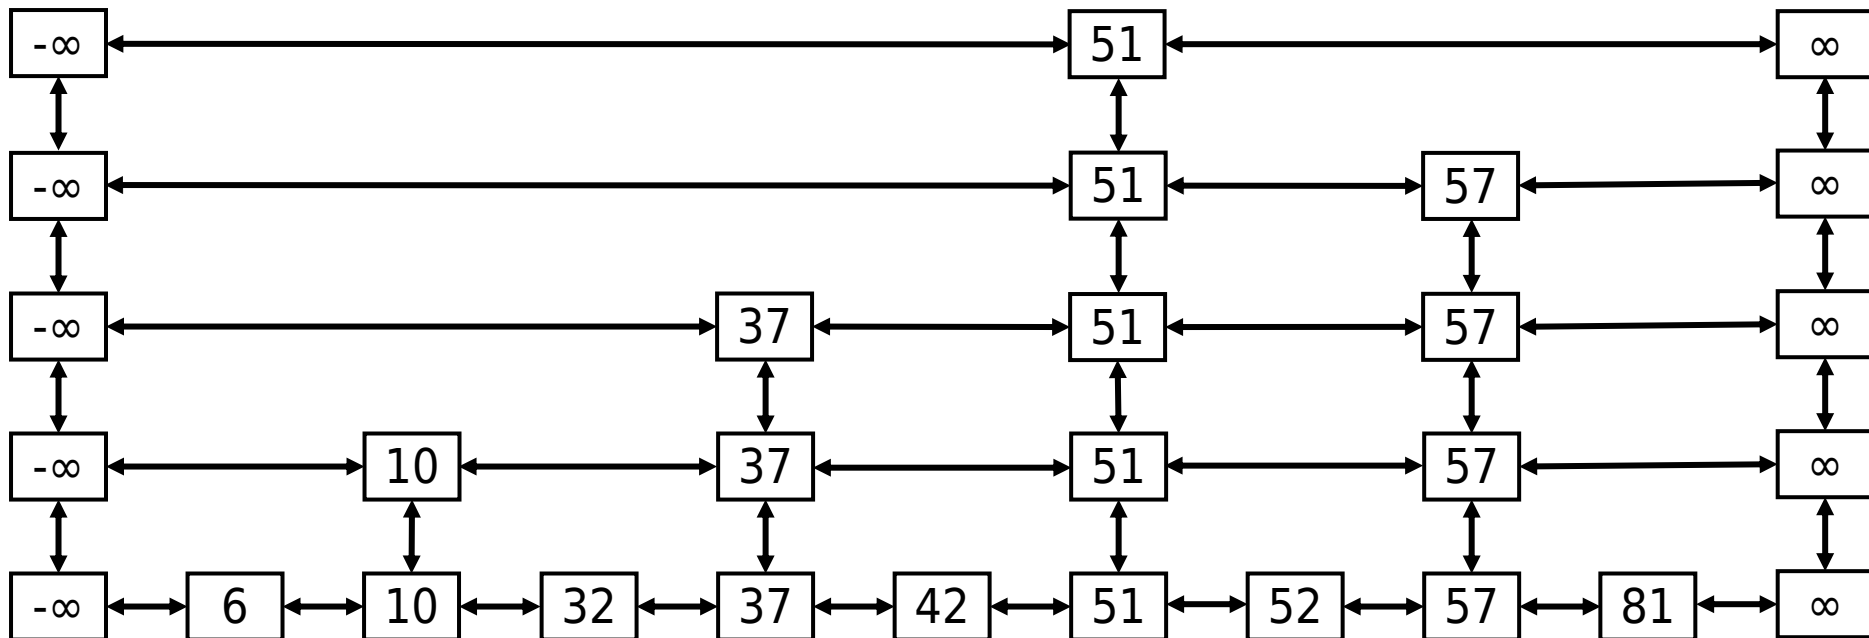
The real **SkipList**

Insert(L,x):

```
level = 1, done = false
while (!done)
  Insert x into level k list.
  Flip a fair coin:
    With probability 1/2: done = true
    With probability 1/2: k = k+1
```

Max level of n -element **SkipList** is $O(\log n)$ with high probability

Example: search 81.

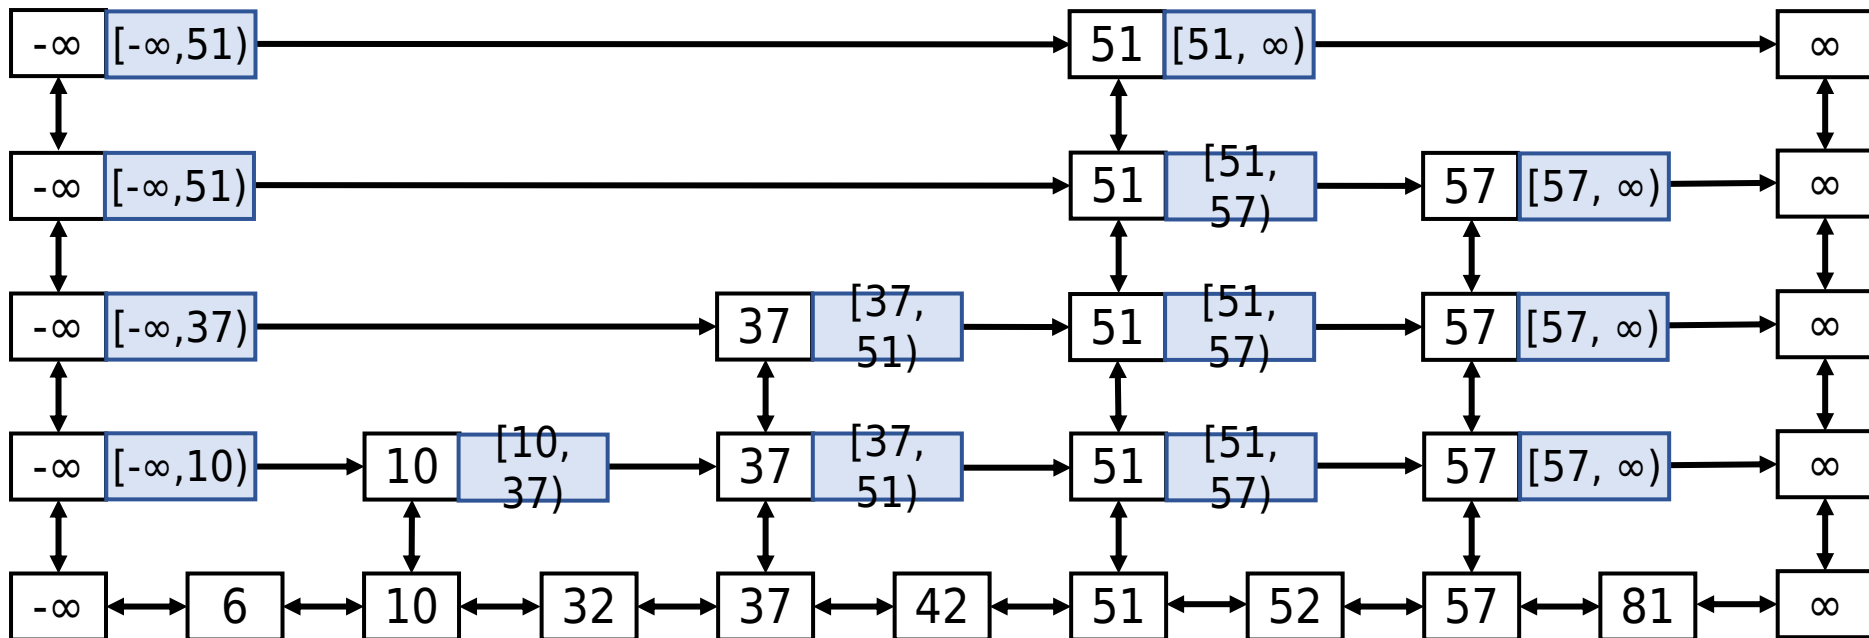


The real **SkipList**

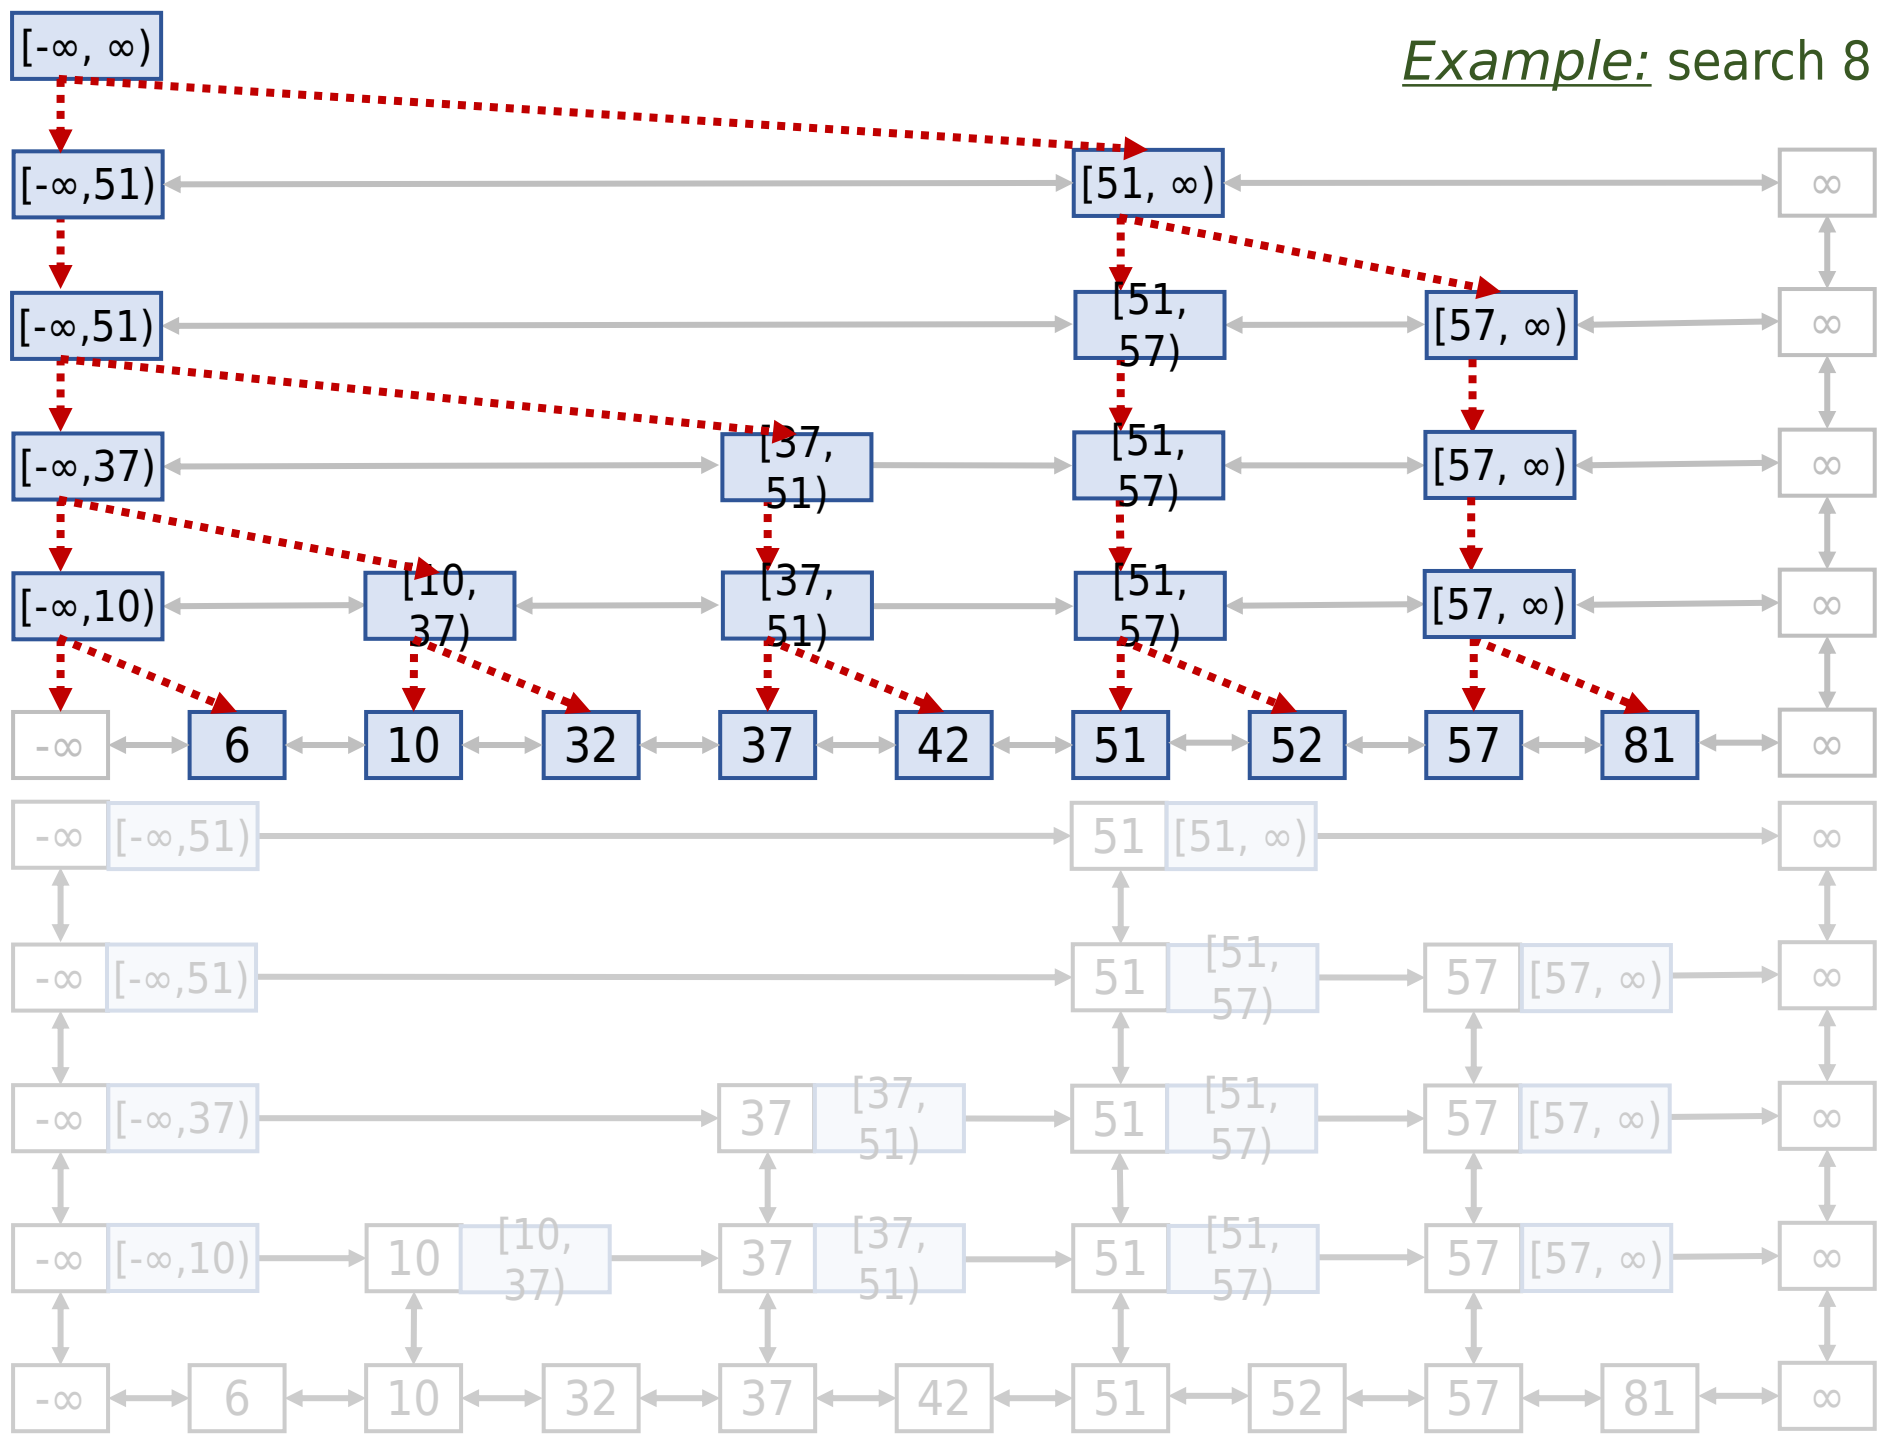
Insert(L,x):

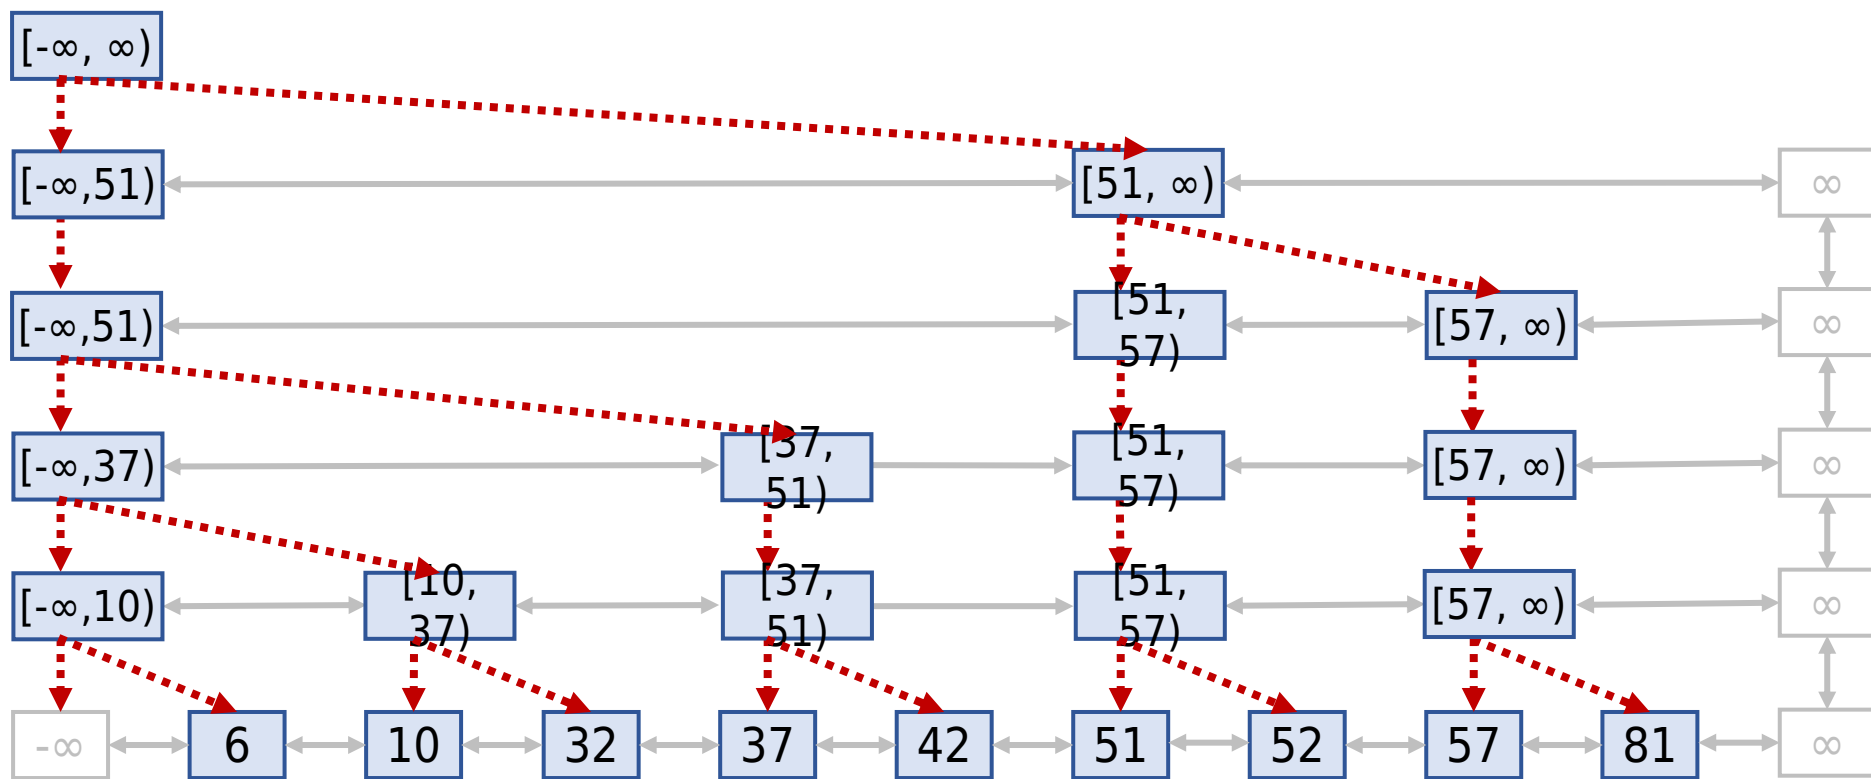
```
level = 1, done = false
while (!done)
  Insert x into level k list.
  Flip a fair coin:
    With probability 1/2: done = true
    With probability 1/2: k = k+1
```

Max level of n -element **SkipList** is $O(\log n)$ with high probability



Example: search 81





Each node in the “search tree” has $O(1)$ children, in expectation.
 Max level of n -element **SkipList** is $O(\log n)$ with high probability.

Let r.v. L be the max level of a n -node SkipList, let C denote the cost for search.

$$\mathbb{P}[L \geq l] \leq n/2^{l-1}$$

$$\mathbb{E}[C] = \mathbb{P}[L \leq \alpha \lg n] \cdot \mathbb{E}[C \mid L \leq \alpha \lg n] +$$

$$\sum_{l=(\alpha \lg n)+1}^{\infty} \mathbb{P}[L = l] \cdot \mathbb{E}[C \mid L = l]$$

Some large constant.

Each node in the “search tree” has $O(1)$ children, in expectation.

Max level of n -element **SkipList** is $O(\log n)$ with high probability.

Let r.v. L be the max level of the skip list, let C denote the cost for search.

$$\mathbb{P}[L \geq l] \leq n/2^{l-1}$$

$$\begin{aligned}\mathbb{E}[C] &= \mathbb{P}[L \leq \alpha \lg n] \cdot \mathbb{E}[C \mid L \leq \alpha \lg n] + \\ &= \sum_{l=(\alpha \lg n)+1}^{\infty} \mathbb{P}[L = l] \cdot \mathbb{E}[C \mid L = l]\end{aligned}$$

$$\leq 1 \cdot O(\lg n) = O(\lg n)$$

$$\begin{aligned}&\leq \sum_{l=(\alpha \lg n)+1}^{\infty} \mathbb{P}[L \geq l] \cdot \mathbb{E}[C \mid L = l] \\ &\leq \sum_{l=(\alpha \lg n)+1}^{\infty} (n/2^{l-1}) \cdot (l + n) \\ &= O(1)\end{aligned}$$

$$\mathbb{E}[C] = O(\log n)$$

That is, search can be done in $O(\log n)$ time in expectation.

Efficient implementation of OSet

	Search(S,k)	Insert(S,x)	Remove(S,x)
BinarySearchTree	$O(h)$ worst-case	$O(h)$ worst-case	$O(h)$ worst-case
Treap	$O(\log n)$ in expectation	$O(\log n)$ in expectation	$O(\log n)$ in expectation
RB-Tree	$O(\log n)$ worst-case	$O(\log n)$ worst-case	$O(\log n)$ worst-case
SkipList	$O(\log n)$ in expectation	$O(\log n)$ in expectation	$O(\log n)$ in expectation

Performance versus Simplicity

Reading

- [CLRS] Ch.13
- [Morin] Ch.4

