

Disjoint Sets

Data Structures and Algorithms

Nanjing University, Fall 2021


郑朝栋

DisjointSet ADT

- A **disjoint-set** ADT maintains a collections $S = \{S_1, S_2, \dots, S_k\}$ of *sets* that are *disjoint* and *dynamic*.
- Each set S_i has a “*representative*” member (i.e., a “*leader*”).
- DisjointSet ADT supports following operations:
 - **MakeSet**(x): create a set containing only x , add the set to S .
 - **Union**(x, y): find the sets containing x and y , say S_x and S_y ;
remove S_x and S_y from S , add $S_x \cup S_y$ to S .
 - **Find**(x): return a pointer to the leader of the set containing x .
- Does not support “remove” elements, or “split” sets.

Sample application of DisjointSet ADT


Computing connected components



Edge processed	Collection of disjoint sets									
MakeSet(\cdot)	$\{a\}$	$\{b\}$	$\{c\}$	$\{d\}$	$\{e\}$	$\{f\}$	$\{g\}$	$\{h\}$	$\{i\}$	$\{j\}$
Union(b,d)	$\{a\}$	$\{b,d\}$	$\{c\}$		$\{e\}$	$\{f\}$	$\{g\}$	$\{h\}$	$\{i\}$	$\{j\}$
(e,g)	$\{a\}$	$\{b,d\}$	$\{c\}$		$\{e,g\}$	$\{f\}$		$\{h\}$	$\{i\}$	$\{j\}$
(a,c)	$\{a,c\}$	$\{b,d\}$			$\{e,g\}$	$\{f\}$		$\{h\}$	$\{i\}$	$\{j\}$
(h,i)	$\{a,c\}$	$\{b,d\}$			$\{e,g\}$	$\{f\}$		$\{h,i\}$		$\{j\}$
(a,b)	$\{a,b,c,d\}$				$\{e,g\}$	$\{f\}$		$\{h,i\}$		$\{j\}$
(e,f)	$\{a,b,c,d\}$				$\{e,f,g\}$			$\{h,i\}$		$\{j\}$
(b,c)	$\{a,b,c,d\}$				$\{e,f,g\}$			$\{h,i\}$		$\{j\}$


Linked-list implementation of DisjointSet

- Basic Idea: Use a linked list to store and represent a set.
- Some more details:
 - A set object has pointers pointing to head and tail of the linked-list.
 - The linked-list contains the elements in the set.
 - Each element has a pointer pointing back to the set object.




Linked-list implementation of DisjointSet

- Basic Idea: Use a linked list to store and represent a set.
- **MakeSet**(x): Create a linked list containing only x . $\Theta(1)$
- **Find**(x): Follow pointer from x back to the set object, then return pointer to the first element in the linked-list. $\Theta(1)$

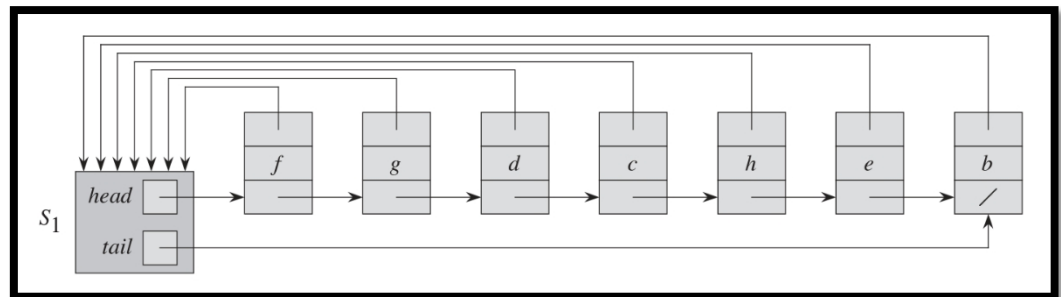


Linked-list implementation of DisjointSet

- Basic Idea: Use a linked list to store elements and represent a set (Set)
 - **Union**(x, y): Append elements of S_y to list in S_x ; destroy set object S_y
- Union can be slow, even in amortized sense!**
- Time depends on size of S_y



Union(g, e)



Linked-list implementation of DisjointSet

- Basic Idea: Use a linked list to store and represent a set.
- **Union**(x, y): Append list in S_y to list in S_x ; destroy set object S_y ;
update set object pointers for elements originally in S_y .

```
MakeSet( $x_0$ )  
for ( $i=1$  to  $n$ )  
  MakeSet( $x_i$ )  
  Union( $x_i, x_0$ )
```

Complexity of this sequence of operations?

$\Theta(n^2)$ in total.

Each **MakeSet** takes $\Theta(1)$ time,
but the average cost of **Union** reaches $\Theta(n)$.

Union operation is too expensive!

Linked-list implementation of DisjointSet

- **Improvement:** Weighted-union heuristic (or, union-by-size).
- **Basic Idea:** In **Union**, append the shorter list to the longer one!

```
MakeSet( $x_0$ )  
for ( $i=1$  to  $n$ )  
  MakeSet( $x_i$ )  
  Union( $x_i, x_0$ )
```

n: Complexity of this sequence of operations?

Worst complexity of any sequence of $n + 1$ MakeSet and then n Union? $O(n \lg n)$

Proof:

- The $n + 1$ **MakeSet** op. take $O(n)$ time in total.
- For **Union** op. whenever its set obj. pointer changes, its set size at least doubles!
- Each element's set obj. pointer changes $O(\lg n)$ times
- The n **Union** op. take $O(n \lg n)$ time in total.

Average cost of **Union** operation is reduced to $O(\lg n)$


Rooted-tree implementation of DisjointSet

- Basic idea: Use a rooted-tree to represent a set; the root of a tree is the “leader” of that set.
- Some details: Each node has a pointer pointing to its parent; parent of a “leader” is the leader itself.
- **MakeSet**(x): Create a tree containing only (root) x ; parent of x is x . $\Theta(1)$
- **Find**(x): Follow parent pointer from x back to the root, and return root.
- **Union**(x, y): Change the parent pointer of the root of y to be the root of x . Time complexity depends on depth of x and y .

Implementation:

- **MakeSet** is fast in both cases.
- **Linked-list**: **Find** is fast, but **Union** is slow.
- **Rooted-tree**: **Find** is slow, but **Union** is fast.


(If **Union** always unions roots of trees.)



Union of

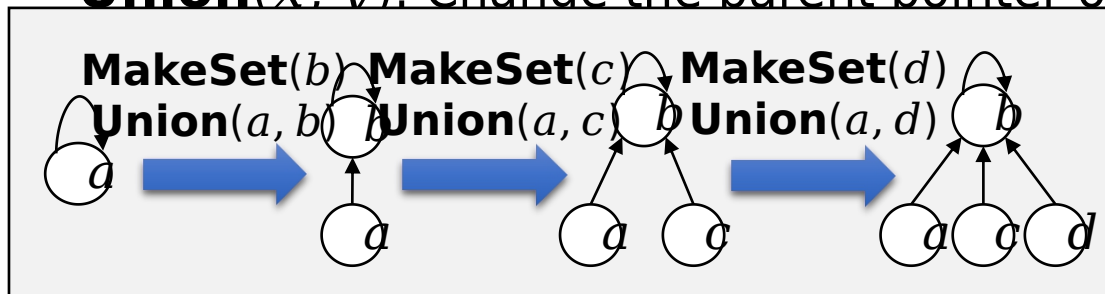
only (root) x ; parent of x

back to the root, and



return root.

- **Union**(x, y): Change the parent pointer of the



```
MakeSet( $x_0$ )
for ( $i=1$  to  $n$ )
  MakeSet( $x_i$ )
  Union( $x_0, x_i$ )
Find( $x_0$ )
```

avg

worst-case cost of **Union** and **Find** to $O(\lg n)$.


(Each time a node's depth increases, the tree size at least doubles. So size n tree has height $O(\lg n)$.)

- Alternatively, use **union-by-height** heuristic:
In **Union**, let tree of larger height.
- Union-by-height reduces worst-case cost of **Union** and **Find** to $O(\lg n)$.

Can we do better?

Rooted-tree implementation of DisjointSet

Path-compression in **Find**




the containing only (root) x ; parent of x

0.

enter from x back to the root, and

height] Change the parent of the root
not of the deep tree. Increase height if

oe
otic
Fin



root r_x ,

make all nodes on this path direct

Find(a)

- Path-compression will not increase

Rooted-tree implementation of DisjointSet

Union-by-height and Path-compression

- ~~**MakeSet**(x): Create a tree containing only (root) x ; parent of x is x .~~
Height of the tree is set to 0.
- **Find**(x): [*path-compression*] Follow parent pointer from x back to root;
let nodes along the path directly point to root; lastly return root.
- **Union**(x, y): [*union-by-height*] Change the parent of the root of the shallow tree to the root of the deep tree. Increase height if necessary.
- **Find** can now change heights! Maintaining heights becomes expensive!
- Maintain *rank*, which is like “height ignoring path-compression.” (rank is always an upper bound of height.)
- **MakeSet**(x): Create a tree containing only (root) x ; parent of x is x .
Rank of the node is set to 0.

Rooted-tree implementation of DisjointSet

Union-by-rank and Path-compression

- **MakeSet**(x): Create a tree containing only (root) x ; parent of x is x .
Rank of the node is set to 0.
- **Find**(x): [*path-compression*] Follow parent pointer from x back to root;
let nodes along the path directly point to root; lastly return root.
- **Union**(x, y): [*union-by-rank*] Change the parent of the root with lower rank to the root with higher rank. Increase rank of new root if necessary.
- Very efficient implementation of DisjointSet, **MakeSet** is $O(1)$, **Find** and **Union** are both *almost $O(1)$ on average*.
- Analysis is highly non-trivial, but we'll do it!

Rooted-tree implementation with union-by-rank and path-compression

Performance Analysis

- **MakeSet**(x): Create a tree containing only (root) x ; parent of x is x .
Rank of the node is set to 0.
- **Find**(x): [*path-compression*] Follow parent pointer from x back to root;
let nodes along the path directly point to root; lastly return root.
- **Union**(x, y): [*union-by-rank*] Change the parent of the root with lower rank to the root with higher rank. Increase rank of new root if necessary.
- **Goal**: Any sequence of **MakeSet**, **Find**, **Union** op. has low avg. cost.
- **Observation**: (a) **MakeSet** can be moved to the beginning of op. sequence, without affecting cost. (b) **MakeSet** itself has low cost.
- **New Goal**: Starting from a forest containing n nodes, any sequence of **Find** and **Union** op. has low avg. cost

Rooted-tree implementation with union-by-rank and path-compression


Performance Analysis

- **MakeSet**(x): Create a tree containing only (root) x ; parent of x is x .
Rank of the node is set to 0.
- **Find**(x): [*path-compression*] Follow parent pointer from x back to root;
let nodes along the path directly point to root; lastly return root.
- **Union**(x, y): [*union-by-rank*] Change the parent of the root with lower rank to the root with higher rank. Increase rank of new root if necessary.
- **Goal**: Starting from a forest containing n nodes, any sequence of **Find** and **Union** op. has low avg. cost.
- **Observation**: $\text{Cost}[\text{Union}(x, y)] = \text{Cost}[\text{Find}(x)] + \text{Cost}[\text{Find}(y)] + O(1)$.
- **New Goal**: Starting from a forest containing n nodes, any sequence of **Find** and **Union** op. has low avg. cost, in which input parameters to **Union** op. are always set leaders

Rooted-tree implementation with union-by-rank and path-compression

Performance Analysis


- **Find**(x): [*path-compression*] Follow parent pointer from x back to root; let nodes along the path directly point to root; lastly return root.
- **PartialFind**(x, y): [*y is ancestor of x*] Follow parent pointer from x back to y ; let nodes along the path point to y 's parent; return parent of y .
- **Goal:** Starting from a forest containing n nodes, any sequence of **Find** and **Union** op. has low avg. cost, in which input parameters to **Union** op. are always set leaders.



Rooted-tree implementation with union-by-rank and path-compression

Performance Analysis

- **PartialFind**(x, y): [*y is ancestor of x*] Follow parent pointer from x back to y ; let nodes along the path point to y 's parent; return parent of y .
- **Union**(x, y): [*union-by-rank*] Change the parent of the root



Rooted-tree implementation with union-by-rank and path-compression


Performance Analysis

- **PartialFind**(x, y): [*y is ancestor of x*] Follow parent pointer from x back to y ; let nodes along the path point to y 's parent; return parent of y .
- **Union**(x, y): [*union-by-rank*] Change the parent of the root with lower rank to the root with higher rank. Increase rank of new root if necessary.
- **Goal:** Starting from a forest containing n nodes, any sequence of **PartialFind** and **Union** op. has low avg. cost, in which every **Union** occurs before any **PartialFind**, and input parameters to **Union** op. are always set leaders.
- **Observation:** Each **Union** op. only costs $O(1)$.
- **New Goal:** Starting from a forest containing n nodes, any sequence of m **PartialFind** op. has low avg. cost.
- **Observation:** Cost of **PartialFind** is dominated by pointer assignments.
- **New Goal:** Starting from a forest containing n nodes

Rooted-tree implementation with union-by-rank and path-compression

Performance Analysis

- **PartialFind**(x, y): [*y is ancestor of x*] Follow parent pointer from x back to y ; let nodes along the path point to y 's parent; return parent of y .
- **Goal:** Starting from a forest containing n nodes, any sequence of m **PartialFind** op. has *low total pointer assignments*.
- $T(m, n, r)$: worst # of ptr. assignments in any seq. of m **PartialFind**, starting from a size n forest where each node has rank at most r .
- **Goal:** $T(m, n, r)$ is small.
- **Claim:** $T(m, n, r) \leq nr$.
- **Proof:** Each node can change parent at most r times, since each new parent has higher rank than the old one.




PartialFind,

starting from a size n forest where each node has rank at most r .

- Fix forest F of n nodes with max rank r , and a seq. C of m **PartialFind** on F .
- $T'(F, C)$: total # of ptr. assignments occurred in C .
- Let s be an arbitrary positive rank, partition F into F_- and F_+ .
- [High Forest] F_+ : containing nodes with rank $> s$;
[Low Forest] F_- : containing nodes with rank $\leq s$.
- Let $|F_+| = n_+$, and $|F_-| = n_-$
- m_+ : number of ops. in C that involve any node in F_+

Rooted-tree implementation with union-by-rank and path-compression



- Consider a **PartialFind**(x, y) in C :
- If $rank(x) > s$: the op. is a **PartialFind** op. in F_+ .
- If $rank(y) \leq s$: the op. is a **PartialFind** op. in F_- .
- If $rank(x) \leq s$ and $rank(y) > s$:
 Split the op. into (a) a **PartialFind** op. in F_+ ; (b) some *shatter* op. in F_- ;
 and (c) a pointer assignment for the “topmost” node in F_+

Consider a **PartialFind**(x, y) in C :

- If $rank(x) > s$: the op. is a **PartialFind** op. in F_+ .
- If $rank(y) \leq s$: the op. is a **PartialFind** op. in F_- .
- If $rank(x) \leq s$ and $rank(y) > s$:
Split the op. into (a) a **PartialFind** op. in F_+ ; (b) some **shatter** op. in F_- ;
and (c) a pointer assignment for the “topmost” node in F_- .

We have converted C into:

- (a) C_+ : ops involving nodes only in F_+ ; (b) C_- : ops involving nodes only in F_- ;
(c) shatter ops; and (d) pointer assignments for “topmost” nodes in F_- .

Observation: each node get shattered at most once (then be “topmost” n

Observation: there are at most m_+ pointer assignments for “topmost” nodes in F_- .

$$T'(F, C) \leq T'(F_+, C_+) + T'(F_-, C_-) + n + m_+$$

Rooted tree implementation with

Any sequence of m **Union** and **Find** on a size n forest --
 takes $O(m + 2n \lg^* n)$ time, even in worst-case.

- $T'(F, C) \leq T'(F_+, C_+) + T'(F_-, C_-) + n + m_+$
- Nodes in F_+ has rank at least $s + 1$ and at most r ;
 Nodes in F_- has rank at most s .
- **Strategy**: obtain a bound of $T'(F_+, C_+)$ to get recurrence of $T'(F, C)$.
- **Claim**: $T(m, n, r) \leq nr$.
 (Recall $T(m, n, r)$: worst # of ptr. assignments in any seq. of m **PartialFind**, starting from a size n forest where each node has rank at most r .)
- **Claim**: There are at most $n/2^i$ nodes of rank i in any size n forest.
- $T'(F_+, C_+) \leq n_+ \cdot r \leq \left(\sum_{i>s} \frac{n}{2^i} \right) \cdot r = \frac{nr}{2^s}$
- Fix $s = \lg r$, then $T'(F, C) \leq T'(F_-, C_-) + 2n + m_+$,
 or equivalently, $T'(F, C) - m \leq (T'(F_-, C_-) - m_-) + 2n$
- $T''(m, n, r) \leq T''(m, n, \lg r) + 2n$, where $T''(m, n, r) = T(m, n, r) - m$
- $T''(m, n, r) \leq 2n \lg^* r$. That is: $T(m, n, r) \leq m + 2n \lg^* r$.

Actual performance is even better!

Summary

- DisjointSet ADT: **MakeSet**(x), **Union**(x, y), and **Find**(x).
- Linked-list based implementation:
 - Use a linked-list to denote a set, first element in list is leader.
 - **Union** is slower, **Find** is fast.
 - With *union-by-size*, **Union** has *average* cost $O(\lg n)$.
- Rooted-tree based implementation:
 - Use a rooted-tree to denote a set, root of the tree is leader.
 - **Union** is fast (if input arg. are leaders), **Find** is slower.
 - With *union-by-size* or *union-by-height*, **Union** and **Find** has *worst-case cost* $O(\lg n)$.
 - With *union-by-rank* and *path-compression*, **Union** and **Find** has *average cost* $O(\lg^* n)$.
(More careful analysis leads to an even better bound!)

Reading

- [CLRS] Ch.21 (excluding 21.4)
- Compared with CLRS, following material presents a simpler analysis for the performance of the DisjointSet data structure when both union-by-rank and path-compression are used.
- [Weiss] Ch.8 (8.6)
- Lecture notes by Jeff Erickson:
<http://jeffe.cs.illinois.edu/teaching/algorithms/notes/11-unionfind.pdf>

