

Ontologies and Databases



What is an Ontology?

A **model** of (some aspect of) the world

- Introduces **vocabulary** relevant to domain
 - Often includes names for classes and relationships
- Specifies intended **meaning** of vocabulary
 - Typically formalised using a suitable logic
 - E.g., OWL formalised using **SHOIQ** description logic
- Consists of two parts
 - Set of **axioms** describing **structure** of the model
 - Set of **facts** describing some particular concrete situation

Axioms

Describe the **structure of the model**, e.g.:

Class: HogwartsStudent

EquivalentTo: Student and attendsSchool
value Hogwarts

Class: HogwartsStudent

SubClassOf: hasPet only (Owl or Cat or Toad)

ObjectProperty: hasPet

Inverses: isPetOf

Class: Phoenix

SubClassOf: isPetOf only Wizard

Facts

Describe some **particular concrete situation**, e.g.:

Individual: Hedwig

Types: Owl

Individual: HarryPotter

Types: HowgwartsStudent

Facts: hasPet Hedwig

Individual: Fawkes

Types: Phoenix

Facts: isPetOf Dumbledore

Obvious Database Analogy

- Ontology axioms analogous to DB **schema**
 - Schema describes structure of and constraints on data
- Ontology facts analogous to DB **data**
 - Instantiates schema
 - Consistent with schema constraints
- But there are also important differences...

Database -v- Ontology

Database:

- Closed world assumption (**CWA**)
 - Missing information treated as false
- Unique name assumption (**UNA**)
 - Each individual has a single, unique name
- Schema behaves as **constraints** on structure of data
 - Define legal database states

Ontology:

- Open world assumption (**OWA**)
 - Missing information treated as unknown
- **No UNA**
 - Individuals may have more than one name
- Ontology axioms behave like **implications** (inference rules)
 - Entail implicit information

Database -v- Ontology

- E.g., given facts/data:

Individual: HarryPotter

Facts: hasFriend RonWeasley

hasFriend HermioneGranger

hasPet Hedwig

Individual: Draco Malfoy

- Query: Is Draco Malfoy a friend of HarryPotter?
 - DB: No
 - Ontology: Don't Know
 - OWA (didn't say Draco was not Harry's friend)

Database -v- Ontology

- E.g., given facts/data:

Individual: HarryPotter

Facts: hasFriend RonWeasley

hasFriend HermioneGranger

hasPet Hedwig

Individual: Draco Malfoy

- Query: How many friends does Harry Potter have?
 - DB: 2
 - Ontology: at least 1
 - No UNA (Ron and Hermione may be 2 names for same person)

Database -v- Ontology

- E.g., given facts/data:

Individual: HarryPotter

Facts: hasFriend RonWeasley
hasFriend HermioneGranger
hasPet Hedwig

Individual: Draco Malfoy

➔ **DifferentIndividuals: RonWeasley HermioneGranger**

- Query: How many friends does Harry Potter have?
 - DB: 2
 - Ontology: at least 2
 - OWA (Harry may have more friends we didn't mention yet)

Database -v- Ontology

- E.g., given facts/data:

Individual: HarryPotter

Facts: hasFriend RonWeasley

hasFriend HermioneGranger

hasPet Hedwig

➔ **Types: hasFriend only RonWeasley or HermioneGranger**

Individual: Draco Malfoy

DifferentIndividuals: RonWeasley HermioneGranger

- Query: How many friends does Harry Potter have?
 - DB: 2
 - Ontology: 2!

Database -v- Ontology

- Insert new facts/data:
 - Individual: Dumbledore
 - Individual: Fawkes
 - Types: Phoenix
 - Facts: isPetOf Dumbledore
- Response from DBMS?
 - Update rejected: **constraint violation**
 - Range of hasPet is Human; Dumbledore is not Human (CWA)
- Response from Ontology reasoner?
 - **Infer** that Dumbledore is Human (range restriction)
 - Also infer that Dumbledore is a Wizard (only a Wizard can have a phoenix as a pet)

DB Query Answering

- Schema plays **no role**
 - Data must explicitly satisfy schema constraints
- Query answering amounts to **model checking**
 - I.e., a “look-up” against the data
- Can be very **efficiently implemented**
 - Worst case complexity is low (logspace) w.r.t. size of data

Ontology Query Answering

- Ontology axioms play a powerful and **crucial role**
 - Answer may include implicitly derived facts
 - Can answer conceptual as well as extensional queries
 - E.g., Can a Muggle have a Phoenix for a pet?
- Query answering amounts to **theorem proving**
 - I.e., logical entailment
- May have very **high worst case complexity**
 - E.g., for OWL, NP-hard w.r.t. size of data (upper bound is an open problem)
 - Implementations may still behave well in typical cases

Ontology Based Information Systems

- Analogous to **relational database management systems**
 - Ontology $\frac{1}{4}$ schema; instances $\frac{1}{4}$ data
- Some important **(dis)advantages**
 - + (Relatively) easy to maintain and update schema
 - Schema plus data are integrated in a logical theory
 - + Query answers reflect both schema and data
 - + Can deal with incomplete information
 - + Able to answer both intensional and extensional queries
 - Semantics may be counter-intuitive or even inappropriate
 - Open -v- closed world; axioms -v- constraints
 - Query answering (logical entailment) much more difficult
 - Can lead to scalability problems

Ontology Based Information Systems

- Similar to **relational**
 - Ontology ¼ schema
- Some important differences
 - + (Relatively) easy to write
 - Both schema and data
 - + Query answering is easier
 - Can answer more complex queries
 - + Able to answer more complex queries
 - Can answer more complex queries
 - Semantics matter
 - Open -v- closed world
 - Query answering is more difficult
 - Can lead to scalability problems



Very powerful, but not miraculous!

Best of Both Worlds?

- W3C OWL working group is developing **OWL 2**
 - OWL 2 is an update to OWL adding many useful features
 - Increased expressive power, e.g., w.r.t. properties
 - Extended support for datatypes and values
 - Database style keys
 - Rich annotations
- OWL 2 also defines several **profiles**
 - Profile is a language subset with
 - Useful computational properties
 - Useful implementation possibilities



Best of Both Worlds?



EL++ profile

- Maximal language for which reasoning (including query answering) known to be worst-case **polynomial**
- Captures expressive power used by many **large-scale ontologies**
 - Features include existential restrictions, intersection, subClass, equivalentClass, class disjointness, range and domain, transitive properties, ...
 - Missing features include value restrictions, Cardinality restrictions (min, max and exact), disjunction and negation

Best of Both Worlds?



DL-Lite profile (not to be confused with OWL Lite!)

- Maximal language for which reasoning (including query answering) is known to be worst case **logspace** (same as DB)
- Captures (most of) expressive power of **ER/UML** schemas
 - Features include limited form of existential restrictions, subClass, equivalentClass, disjointness, range and domain, symmetric properties, ...
- Query answering can be implemented using **query rewriting**
 - Resulting SQL query/queries capture all information from axioms
 - Can use query/queries with standard DBMS and relational data

Best of Both Worlds?



OWL-R profile

- Allows for scalable (**polynomial**) reasoning using rule-based technologies
- Includes support for **most OWL features**
 - But standard semantics only apply when they are used in a restricted way
 - Related to DLP and pD*
- Can be implemented on top of **rule extended DBMS**
 - E.g., Oracle's OWL Prime implemented using forward chaining rules in Oracle 11g

Summary

- Ontologies consist of sets of axioms and facts
- Analogous to DB: axioms $\frac{1}{4}$ schema; facts $\frac{1}{4}$ data
- Important differences in semantics
 - DB: UNA, CWA and constraints
 - Ontology: OWA and implications
- Ontologies are very powerful, but there are costs
 - Can be scalability problems
- OWL 2 provides choice of several profiles
 - Tractable reasoning (logspace or polynomial)
 - Different features and implementation pathways